A parameter-sparse state-space framework for characterizing coastal dune-aquifer architecture from manual dipwell records

M. Hollingham

Draft — 2026

Abstract

Study region: Newborough Warren, Anglesey, Wales. This is a coastal dune aquifer on the Irish Sea, monitored by 88 manual dipwells over 21 years and one climate station.

Study focus: Dune-slack vegetation depends on water-table depth within a band a few tens of centimetres wide. Distributed flow models of dune sand, however, need conductivity and thickness fields that are costly and poorly constrained. We ask whether routine dipwell records alone can resolve dune-aquifer architecture. We couple hierarchical clustering of hydrograph dynamics to a three-term state-space model, with terms for recharge, atmospheric draw and head-dependent drainage above a fixed datum. The model is fitted at each well from monthly rainfall and temperature, without coring, hydraulic testing or flow-model calibration.

New hydrological insights for the region: Five hydrogeological clusters emerge, each with distinct mechanistic coefficients. Interception-corrected specific yields converge across them, and the coastal forest cluster sits at the upper edge of the open-dune range. The behavioural contrasts are therefore consistent with surface and land-cover control rather than storage. The fit closes the mean balance to within 1.8% of losses, although its drainage–evapotranspiration split depends on the datum. It raises median iterative Nash–Sutcliffe efficiency from 0.42 to 0.77 over a linear benchmark. A coastal-retreat gradient is separated from the site-wide climate background; it decays over an inland reach of about 900 m to 1 km. It accounts for about half of the decline of the most coast-proximal cluster, and for most of the smaller decline of the western dune cluster.

Keywords: coastal dune aquifer; dune slack; state-space model; specific yield; groundwater characterization; coastal retreat

Highlights

- Dune-aquifer architecture resolved from manual dipwells and one climate station
- Coefficients and water balance reported at a declared datum, chosen by a published sweep
- Iterative forecast skill rises from NSE 0.42 to 0.77 over a linear benchmark

- Interception-corrected specific yields converge across clusters: storage is broadly uniform
- Coastal retreat resolved as a forcing distinct from climate on long-term decline

Plain-language summary

Dune slacks, the damp hollows between sand dunes, hold some of the most water-sensitive habitats in north-west Europe. Their plants each live within a band of summer water depth only a few tens of centimetres wide. At most sites the only record of that water is a set of pipes read by hand once a month. This paper asks whether such a record, with one nearby weather station, is enough to work out how a dune aquifer is built and how it behaves, without drilling or a detailed model of underground flow.

At Newborough Warren on Anglesey, 88 such pipes have been read for 21 years. Sorted by the shape of their records, they fall into five groups: a quick-responding eastern block beside the lake and on the open dune, a slower western block, and two groups under the pine plantation. A simple monthly model, in which rain raises the water table and evaporation and drainage lower it, describes each group with three numbers. Run month after month from the weather alone, it follows the water table far better than the same model without drainage: a typical score of 0.77 against 0.42, on a scale where 1 is perfect.

The amount of water the sand holds turns out to be much the same everywhere once the rain caught by the forest canopy is allowed for. The forest groups behave differently mainly because the trees intercept rain at the surface, not because the ground beneath them is different. How the water lost each year divides between drainage and evaporation depends on an assumption about the depth the water drains towards, and it is reported with that condition.

The western shore is eroding, and the water table near it is falling faster than inland: by roughly 26 mm a year at 150 m from the shore, fading over about a kilometre. This coastal effect accounts for about half of the summer decline nearest the coast. It has to be separated out before any trend in the water table is put down to climate.

The method needs only routine records, so other dune sites can use it. What it cannot show is how water moves from one part of a site to another; that needs measurements of the sand itself.

1. Introduction

Coastal dune slacks are among the most hydrologically sensitive habitats in north-west Europe. They are defined by a seasonally fluctuating water table, and that water table sits within a narrow envelope. The wet- and dry-slack vegetation communities of conservation interest occupy a band of summer-minimum depth only a few tens of centimetres wide (Davy et al., 2006; Curreli et al., 2013). The envelope is set by groundwater, not by surface management. The conservation status of a dune system therefore cannot be understood without a

quantitative account of how its aquifer fills, stores and releases water. That account must cover how recharge enters the water table, how evaporative demand and lateral drainage remove it, and how these processes vary across the geomorphological range of the site.

Characterizing that architecture by conventional hydrogeological means is difficult in dune settings. Distributed groundwater models of the MODFLOW family require hydraulic conductivity fields, aquifer-thickness surfaces and boundary fluxes. These are expensive to measure and poorly constrained in unconsolidated, rapidly transmissive dune sand. At Newborough Warren itself, an earlier distributed-model effort encountered exactly these calibration difficulties (Betson et al., 2002). Slug and pumping tests sample only the immediate vicinity of a borehole. Interpolating a parameter field across a dune system of several square kilometres needs dense borehole networks, and these are rarely available.

Yet many dune sites are monitored, often over decades, by exactly the kind of low-cost instrumentation that distributed models cannot directly exploit: networks of manual dipwells read at monthly intervals, alongside a nearby meteorological station. This paper addresses a methodological question. Is that routine monitoring record, on its own, sufficient to resolve the hydrogeological architecture of a dune aquifer — its storage, recharge sensitivity and drainage behaviour? And can it do so with enough spatial detail to support conservation interpretation, without recourse to coring, hydraulic testing or calibrated continuous-flow modelling?

We develop and demonstrate such a framework. Its core is a three-term state-space model (SSM) fitted independently at each well. The model partitions the monthly change in water-table position into three terms: a recharge term proportional to rainfall, an atmospheric-draw term proportional to potential evapotranspiration, and a head-dependent drainage term proportional to the water-table displacement above a fixed reference datum. Drainage is referenced to a datum below ground rather than to the ground surface. This captures the head-dependence of lateral discharge without calibrating the drainage geometry. It also resolves a positive drainage coefficient at every cluster, including the low-gradient forest interior, where a drainage term referenced to the ground surface is not identified. A companion methods paper (Hollingham, 2026c) sets out the model, the choice of datum and what the no-intercept form identifies; here the model is used as a characterization tool. The three fitted coefficients are physically interpretable: recharge sensitivity, atmospheric draw and proportional drainage decay. As forcing, they require only monthly rainfall and temperature from a single station.

Around this core we assemble a set of complementary, independently derived characterizations that cross-validate one another. First, hierarchical clustering of hydrograph *dynamics* partitions the network into spatially coherent hydrogeological units before any model is fitted. It uses a correlation-based distance that isolates temporal behaviour from absolute water-table elevation. This provides a structure that the mechanistic coefficients must then independently reproduce. A pattern-matching affinity analysis extends the classification to wells with shorter records and identifies the gradational transition zones between units.

Next, the SSM is benchmarked against a traditional linear model under iterative simulation. The spatial pattern of the predictive gain itself then becomes a non-invasive diagnostic of where head-dependent drainage governs the aquifer. A water-balance decomposition built directly

from the fitted coefficients quantifies the recharge–loss partition at each cluster under identical climate forcing. An independent water-table-fluctuation estimate of specific yield tests whether the behavioural contrasts between units reflect genuine differences in subsurface storage, or boundary conditions imposed at the surface. Finally, interpolation of the per-well coefficients yields continuous maps of recharge sensitivity, atmospheric draw and drainage rate. It also yields a mean head surface with its associated flow field and, at each well, a water-balance residual that makes explicit where the three-term model closes.

We also resolve an external forcing that conventional climate-trend analysis at a dune site can misattribute. Where a shoreline is actively eroding, the seaward drainage boundary migrates progressively landward. This deepens the water table near the coast independently of climate, and so imposes a spatial gradient on long-term summer-minimum trends. We fit a physics-based, network-scale regression of water-table trend against distance from the eroding shoreline. We use it to partition each cluster's summer-minimum decline into a coastal-retreat component and a climate-driven background. The distinction matters both for interpreting the aquifer's long-term trajectory and for any subsequent attribution of trends to climate alone.

The framework is demonstrated at Newborough Warren, a coastal dune system on Anglesey, north-west Wales. The site has been monitored over twenty-one years by a network of eighty-eight dipwells (eighty-nine monitoring points, including the Llyn Rhos-Ddu lake gauge). The network spans the full geomorphological range of the site, from lake-margin and open-dune slacks to a mature coniferous plantation on the elevated landward flank. The density and length of this record make it an unusually well-suited proving ground. It is large enough to test whether the framework recovers structure that is independently visible in the raw hydrographs. It is also varied enough to span the contrast between a shallow, fast-cycling open dune and a slower-draining forested interior.

The objectives of this paper are therefore: (i) to apply the displacement-form state-space model of Hollingham (2026c) as a characterization framework; (ii) to demonstrate that hydrograph-dynamics clustering, mechanistic SSM coefficients, model-benchmarking diagnostics, water-balance decomposition and specific-yield estimation converge on a single, internally consistent picture of aquifer architecture; (iii) to map that architecture continuously across the site and identify where the three-term water balance closes and where it does not; (iv) to quantify the coastal-retreat forcing on long-term water-table trends; and (v) to assess the framework's transferability to other dune systems monitored by routine dipwell networks.

Two management interventions were undertaken at the site: dune scraping and a partial plantation clearfell. Their hydrological consequences are quantified separately within a before-after-control-impact design and are reported in a companion paper (Hollingham, 2026b). They are not treated here; the plantation enters only as a land-cover boundary condition on the aquifer characterization. The intervention dates — the December 2017 clearfell and the dune-scrape years — are the author's recollection, from twenty-one years monitoring the site. The April 2015 CEH36 scrape is separately documented. No documentary record of the operations has been obtained.

2. Study site

Newborough Warren occupies the south-western tip of the Isle of Anglesey in north-west Wales (Figure 1). The modern dune system is a late-Holocene feature. The present dunes developed over previously cleared agricultural land following storms in 1330 (Ranwell, 1958; Rhind et al., 2001). The site has no perennial surface drainage. Flow lines shown on later figures are DEM-derived flow-accumulation paths, included for spatial reference.

Weakly permeable glacial till underlies the dune field and forms the effective base of the sand aquifer. A bedrock ridge, demarked by the 20 m AOD contour, trends south-south-west to north-north-east. It forms the northern boundary of the study area and outcrops within the northern plantation. The hydrological study area is defined by the extent of the DEM-derived flow network and covers approximately 865 ha. The dipwell network spans a core zone of roughly 720 ha within it, lying entirely on the south-eastern flank of the ridge. Dune terrain on the north-western flank drains toward the Malltraeth estuary. It is hydrologically separate and is not considered here.

The substrate divides the aquifer south-east of the ridge into two contrasting blocks (Stratford et al., 2007). In the east, shallow sand over till and estuarine deposits constrains storage and produces a responsive, flashy water-table regime. In the west, deep aeolian sand has accumulated over a broader platform. This permits substantial groundwater storage and strongly buffered seasonal fluctuations. The present data do not resolve what causes this east–west contrast (Section 5.1). It may reflect two distinct substrate bodies. Alternatively, it may reflect a single estuarine-clay base that outcrops on the higher eastern and southern ground and dips beneath the surface toward Caernarfon Bay, thinning the sand column in the east and thickening it toward the south-west.

The aquifer discharges principally across the south-western foreshore of Caernarfon Bay and the southern and eastern coast of the Menai Strait. Llyn Rhos-Ddu, a small lake within the warren, acts as a near-fixed-head body. It is the dominant local drainage sink for the immediately adjacent wells. It receives groundwater throughout the year, because the regional water table maintains its head above the lake surface (Stratford et al., 2007).

The northern part of the study area is managed plantation dominated by mature Corsican pine. The plantation enters the present analysis solely as a land-cover boundary condition. Its canopy intercepts a fraction of incident rainfall before it reaches the ground. This interception is treated below as a modifier of the surface input to the forested clusters, not as a management variable. The dune-slack communities of the open system are classified under the National Vegetation Classification as SD15b (wet slack) and SD16 (dry slack). Their community-mean water-table levels (Curreli et al., 2013) provide the ecological reference frame for the site, but are not used analytically in this paper.

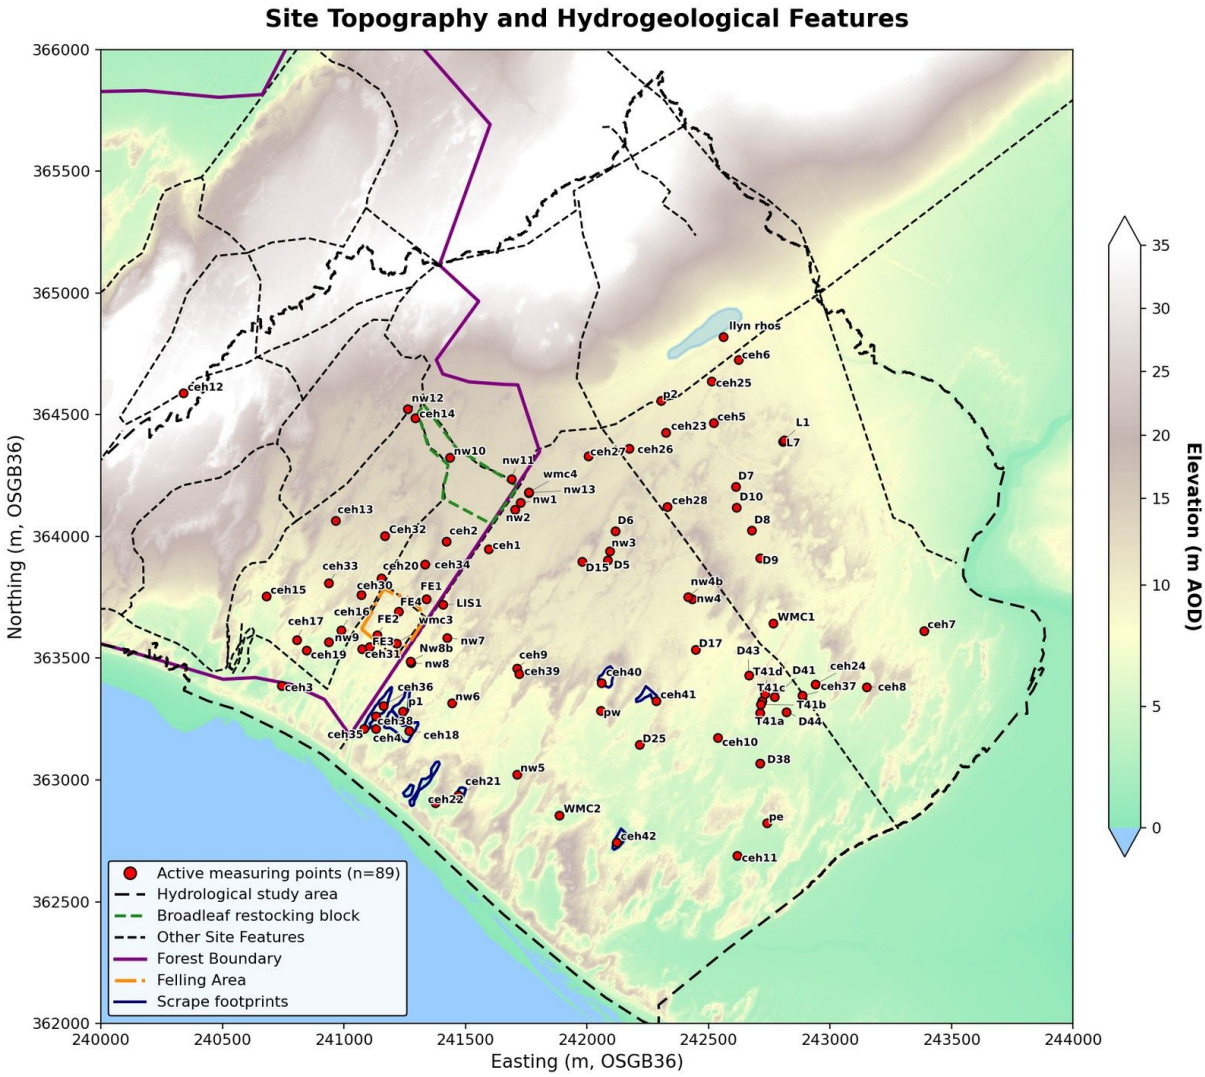


Figure 1. *Site topography, geology and monitoring network at Newborough Warren. (Digital elevation model with rock ridge, forest boundary and watershed boundary. No perennial surface streams are present on site. LiDAR data: © NRW & OS.)*

3. Data and methods

3.1 Groundwater and climate records

The primary observational dataset is a network of manually read dipwells. Water-table elevation is recorded at approximately monthly intervals. Each reading is taken at the end of the month and represents the water level for the month just ended. Readings are bucketed to a monthly timestamp by field convention: a reading on or before the fifteenth of a month is assigned to the preceding month, and a later reading to the same month. All monthly series therefore index the month to which the reading belongs, not the day on which it was taken. Of the full monitoring array, wells with sufficient record length form the reference network used

for model fitting. The remaining wells, with shorter records, form an extended network classified by pattern matching (Section 3.3). Water levels are expressed both in metres above ordnance datum and, for the displacement formulation, as depth below the ground surface at each well.

Climate forcing is taken from the Royal Air Force Valley station. Monthly precipitation totals are used directly. Potential evapotranspiration (PET) is computed by the Thornthwaite method from monthly mean temperature and station latitude. The heat index is summed over the trailing twelve months rather than the calendar year, so that it is defined for a record that does not end in December. Thornthwaite PET is an empirical index of atmospheric evaporative demand. Unlike Penman–Monteith, it does not define a specific reference surface, and this property bears on the interpretation of the atmospheric-draw coefficient (Section 5.3). Rainfall and PET are monthly totals for the same calendar month as the water-table change they are aligned against. No rainfall lag is applied, consistent with the end-of-month measurement convention.

3.2 Hierarchical clustering of hydrograph dynamics

The reference network is partitioned into hydrogeologically distinct zones by agglomerative hierarchical clustering with Ward's minimum-variance linkage (Ward, 1963). The distance metric is the Pearson correlation distance between standardized well hydrographs. This metric compares the *shape* of the seasonal and inter-annual water-table dynamics and discards absolute elevation. Two wells that fluctuate in synchrony are treated as similar, whatever their mean water-table depth. The clustering therefore isolates the dynamic hydrogeological signature — the timing and relative amplitude of filling and draining — from the topographic offset between wells.

The number of clusters is selected by a dual-metric approach that combines two measures. The first is the Ward's merge-distance profile, which identifies the partition at which further agglomeration incurs a disproportionate increase in within-cluster variance. The second is the mean silhouette coefficient (Rousseeuw, 1987), which favours partitions with well-separated, internally coherent groups. Seasonal amplitude is the long-term mean difference between winter maximum and summer minimum at each cluster centroid. It is retained as a secondary descriptor characterizing each unit, but it is not used in the clustering itself.

3.3 Network extension by Pearson affinity analysis

The primary classification, derived from the reference network, is extended to the wells with shorter records by Pearson affinity analysis. Unlike the mechanistic models, this stage operates on the raw observed time series, without model fitting. Each well's hydrograph is correlated against every cluster centroid, and the well is assigned to the cluster with which it correlates most strongly. The correlation delta between the assigned cluster and the next-strongest one classifies membership into tiers. The tiers distinguish core members, which have an unambiguous dominant affinity, from wells whose affinities to neighbouring clusters compete closely. A multi-cluster affinity flag identifies wells whose hydrographs are nearly equally well matched by several centroids. The spatial concentration of these flags delineates the

gradational transition zones, where the boundaries between hydrogeological units are diffuse rather than sharp. The extended-network wells are classified for spatial completeness and are not used in the cluster-level mechanistic fits.

3.4 State-space model and the displacement formulation

The mechanistic core of the framework is the displacement-form state-space model (SSM) of Hollingham (2026c). It relates the monthly change in water-table position to climate forcing and to the antecedent head. For a given well or cluster centroid, the model is

$$\Delta h(t) = \beta_1 \cdot P(t) - \beta_2 \cdot PET(t) - \beta_3 \cdot h_disp(t-1)$$

where $\Delta h(t)$ is the change in water-table position during month t, and P(t) and PET(t) are the month-t precipitation and potential-evapotranspiration totals. The three coefficients are reported as positive quantities. β_1 (recharge sensitivity) is the water-table rise per unit rainfall. β_2 (atmospheric draw) is the decline per unit PET. β_3 (proportional drainage decay) is the fraction of the displacement above the drainage datum that drainage removes each month.

The drainage term is driven by the displacement above the datum at the start of the month. This avoids regressing a change on its own end-point:

$$h_disp(t-1) = z_0 + h(t-1)$$

where h is the water-table position (negative below ground) and z_0 is the depth of the drainage datum below the ground surface. The model is fitted by ordinary least squares without an intercept. This is the form published here, with one datum at every well: z_0 = 3.7 m below ground. The datum is a modelling choice rather than a measured base. Because no intercept is fitted, the datum is not inert: a change of datum is redistributed across β_1, β_2 and β_3. The coefficients, and the half-life derived from β_3 (Section 4.7), are therefore reported conditional on 3.7 m. Hollingham (2026c) sweeps the datum at this site and sets out what it decides. From 1.6 m downward, every cluster's β_3 is positive and significant. At 3.7 m the mean drainage flux is close to its deepest-datum value. The flux and the drainage share of the mean loss are therefore insensitive to the datum, while β_3 continues to scale with it. The ranking of the clusters by β_3 is the same at every datum swept. A depth below ground fits better than any common elevation. The same paper sets out what the free-intercept form identifies instead, and why the no-intercept form is the one used for the coefficients and the water balance.

The recharge sensitivity of each unit is summarized by the lumped catchment storage coefficient, LCSC = $100/\beta_1$. This is the depth of rainfall required to raise the water table by 100 mm. A lower LCSC denotes greater recharge sensitivity. The fitted coefficients fold together vegetation effects, specific yield and soil-moisture dynamics, and no intercept is fitted. For both reasons, the recharge term in an open-boundary system of this kind can absorb unmeasured boundary fluxes into β_1 (Young, 2011; Knotters and van Walsum, 1997). This is addressed through the residual analysis (Section 3.10) and the specific-yield estimation (Section 3.8).

A depth-dependent extension was evaluated at cluster level as a diagnostic of depth coupling (Section 5.2). In it, β_2 is replaced by $\beta_2 \cdot \exp(-\kappa d_w)$, so that atmospheric draw declines as the

water-table depth d_w increases, and κ is found by grid search. The fixed-β_2 model is used throughout, as an effective depth-averaged parameterization.

The model is fitted at two scales, each on a declared record. The mechanistic characterization (Table 1) is fitted to each cluster-mean hydrograph on its full record. Averaging suppresses single-well noise and returns one coefficient set per zone. A cluster coefficient is therefore a centroid rather than a property of every member, and it depends on the partition. The per-well coefficients are fitted on the most recent 100 months at every well, so that wells are compared under the same climate. These are the coefficients behind the coefficient ranges (Table 5), the coefficient surfaces (Figures 11–13) and the benchmark (Section 3.5). The per-well half-lives (Table 6) are instead fitted on each well's full record, because the slow forest recession is not resolved within 100 months. At the cluster centroids, fitting the 100-month window instead of the full record shifts β_3 by between −3% and +16% outside the Main Forest, and by −34% at C4 (Hollingham, 2026c).

3.5 Model benchmarking against a traditional linear model

The predictive value of the head-dependent drainage term is assessed against a traditional linear model (TLM). The TLM is the same model without the drainage term and with a free intercept. It is fitted at every reference well on the 100-month comparison window and run freely over the same months from the first observed level, so the test is in-sample. Hollingham (2026c) sets out the models, the protocol and the network results, both for the no-intercept form used here and for the free-intercept form. The benchmark tests the drainage term against its absence. It does not choose between the two forms of the SSM, which are compared out of sample in the same paper. Here, the per-well gain in iterative Nash–Sutcliffe efficiency of the SSM over the TLM is mapped across the network (Figure 9). The map is a diagnostic of where the drainage term is indispensable to prediction and where the aquifer approximates a linear climate response.

3.6 Spatial coefficient mapping and within-forest regression

The per-well SSM coefficients are interpolated across the site by inverse-distance weighting. The interpolation runs over a regular grid built on a LiDAR-derived digital elevation model, with ridge cells and non-aquifer areas masked. The result is a set of continuous surfaces of recharge sensitivity, atmospheric draw and drainage rate, together with a model-fit-quality surface. These surfaces give a spatially resolved expression of aquifer behaviour that complements the cluster-level summaries.

Within the forested zone the per-well coefficients span much of the site-wide range. There, they are regressed against elevation, distance from the ridge and easting, and against canopy density measured at each well from the aerial imagery (Section 4.8). The regression tests whether the within-forest variation tracks topographic position and water-table depth, or canopy differences. The subsurface beneath the plantation is not cored, so elevation and position serve as proxies for substrate. The regression thus isolates this topographic control on the mechanistic coefficients beneath a single land cover.

3.7 Mean monthly water-balance decomposition

The fitted cluster coefficients are used to decompose the mean monthly water balance of each cluster into its component fluxes, under identical climate forcing. Recharge is computed as β_1 times mean monthly precipitation. Atmospheric draw is β_2 times mean monthly PET. Gravity drainage is β_3 times the mean displacement above the datum. All terms are expressed in metres of water-table change per month, the natural unit of the state-space model. They are directly comparable across clusters that receive the same climate input. Canopy interception is not added as a separate loss term. In the forested clusters it is already embedded implicitly in the lower fitted β_1 and β_2, and adding it explicitly would double-count losses. The net monthly balance is recharge minus total losses. It quantifies the residual that would be required from an unmeasured boundary input to maintain long-run equilibrium.

An indicative volumetric conversion is presented alongside the head-space decomposition. It uses assumed specific-yield values appropriate to the substrate of each cluster. The empirical specific-yield estimates of Section 3.8 are used to bracket the volumetric fluxes from above. For the forested clusters, canopy interception is applied as a reduction of the surface input, following the site-specific interception fraction (Freeman, 2008). That fraction is a single-plot, single-season measurement, and the only site-specific value available. Interception therefore reduces the boundary input rather than being additive to losses. The evapotranspiration–drainage partition is bracketed by two independent methods: the SSM headspace ratio and a seasonal recession-curve analysis. It is reported as a mid-point with an associated uncertainty range.

3.8 Specific-yield estimation by the water-table-fluctuation method

An independent estimate of specific yield is derived from the dipwell record by the water-table-fluctuation method (Healy and Cook, 2002). In this method, specific yield is the ratio of net recharge to the corresponding water-table rise. The method is applied to the cluster-mean hydrographs. It retains only months in which both the water table rose and net recharge was positive above defined thresholds, and it constrains event-level estimates to a physically plausible range. Monthly data cannot isolate individual storm events. At monthly resolution, the method is also known to conflate gravity drainage with capillary-fringe release. The resulting estimates are therefore interpreted as upper bounds on effective storage rather than as point estimates (Healy and Cook, 2002; Scanlon et al., 2002).

For the forested clusters, a canopy-interception correction is applied to the precipitation term before net recharge is computed. The correction reduces only the rainfall input and leaves the Thornthwaite PET term unchanged. PET is an energy-based atmospheric-demand index, independent of land cover. Reducing it to account for canopy-evaporated interception would conflate atmospheric demand with realized evapotranspiration, and would remove the same quantity from both sides of the balance. The corrected forest estimates are compared against the open-dune range. The comparison tests whether the attenuated water-table response beneath the canopy reflects the surface boundary condition or a genuinely different subsurface storage. The per-well event-median specific yields are interpolated across the site. This assesses

whether the spatial distribution of storage is uniform or shows discontinuities coinciding with the cluster margins.

3.9 Network-scale coastal-retreat gradient regression

Where a shoreline is actively eroding, the landward migration of the seaward drainage boundary acts as a slowly advancing line sink. It lowers the ground surface relative to the water table, which raises the displacement head at adjacent wells and steepens the seaward hydraulic gradient. The near-coast water table therefore deepens over time, independently of climate, along a Dupuit drawdown curve (Fetter, 2001). This mechanism, in which an eroding dune margin draws down its own hinterland, is documented at the retreating Whiteford Burrows dune system in South Wales (Robins, Pye and Wallace, 2013). It imposes a spatial gradient on long-term summer-minimum trends. The aim is to partition each cluster's observed summer-minimum decline into a coastal-retreat component and a climate-driven background. To do so, a network-scale, physics-based non-linear regression is fitted to per-well water-table trends against distance from the eroding shoreline. The perpendicular distance from each well to the eroding Caernarfon Bay high-water mark is computed within the pipeline. It is the minimum distance to a committed coastline polyline (OpenStreetMap mean-high-water, EPSG:27700), restricted to the west-facing Newborough shoreline. The Menai Strait is excluded as a non-eroding boundary.

Two functional forms are fitted. The first is a Dupuit–Forchheimer steady-state strip-aquifer form, in which the coastal anomaly declines linearly to a far-field background and is capped at zero. The second is an exponential-decay form, representing a diffusive transient response. These are the steady-state and transient limits of the same drawdown physics: a Dupuit drawdown curve from the migrating boundary. Its inland reach L_cg is set by the aquifer's hydraulic diffusivity D = Kb/Sy, the same length-scale physics that governs the interior head-loss reach of the forest block. Both forms are fitted by profile non-linear least squares, with well and month fixed effects and a cumulative-water-balance covariate absorbed by within-well demeaning. The Akaike information criterion selects between the two forms (Akaike, 1974).

Three nested specifications are fitted to test robustness to forest-cover inclusion. The first is the full network, excluding the clearfell-zone wells. The second is a forest-free network that drops all forested wells (C4 Main Forest and C5 Coastal Forest). The third is a single-cluster specification that uses only the western non-forested cluster (C3 Western Residual). In this specification the far-field background is held to the network value, because the cluster contains no well at the coast and an unconstrained fit on its restricted distance range would be under-identified. The forest-free linear-with-cutoff fit is taken as the headline. A fourth fit controls forest cover explicitly rather than by exclusion: it refits the full network with a canopy × time term. The coast-edge amplitude is essentially unchanged, and the absorbed canopy effect is a separate −9.0 ± 2.5 mm yr^{-1} deepening under pine. This indicates that forest cover does not drive the gradient. The headline fit is then applied to decompose each cluster's (C1–C5) summer-minimum slope into gradient and climate components. This quantifies the contribution of coastal retreat to long-term water-table trends across the network. The coastal term of each cluster is the mean, over its wells, of the capped gradient, each evaluated at the well's own

distance from the shore. A cluster that straddles the inland reach therefore carries the share of its wells that lie inside it.

3.10 Spatial interpolation, flow field and water-balance residual

The mean water-table surface is built by ordinary kriging of each well's mean head with a spherical variogram. Kriging is the geostatistical approach generally preferred over deterministic interpolators, such as inverse-distance weighting, for mapping groundwater levels from sparse wells (Varouchakis and Hristopulos, 2013; Rao et al., 2022). Kriging admits a drift on the ground surface (Desbarats et al., 2002), but here the drift was tested rather than assumed. Three drifts were compared with ordinary kriging by leave-one-well-out error and by a physical test against Sentinel-2 wet-floor extents, and ordinary kriging was selected (Hollingham, 2026d, §3.8). The edges are anchored rather than extrapolated. Seaward, the anchor is the sea along the high-water mark, at a coastal head fitted to the coastal wells between mean tide level and mean high water (UKHO, 2026). Inland, the anchors are the Afon Braint, Llyn Rhos-Ddu and the ridge well CEH12. No elevation mask is applied. Flow arrows are the negative head gradient, one per 100 m square, sized by the gradient. No hydraulic conductivity or aquifer thickness enters them. They describe the flow geometry of the mapped surface rather than the output of a calibrated continuous-flow model (Bear, 1972; Freeze and Cherry, 1979).

The water-balance residual quantifies, at each reference well, the degree to which the three-term SSM closes the long-term water balance. It is computed from the fitted coefficients as the difference between modelled mean losses (atmospheric draw plus drainage) and modelled mean recharge. Gross rainfall is used at all wells, including the forested ones. Interception is already carried in $\beta_2 \cdot \bar{PET}$, since the SSM is fitted on gross rainfall and above-canopy Thornthwaite PET. Subtracting it from rainfall as well would double-count it. Positive residuals indicate locations where modelled losses exceed modelled recharge. This points to an inadequacy in the three-term model at that location, or in principle to an unmodelled input. Near-zero residuals indicate that the three-term balance closes adequately. The residual is shown at the wells alone and is not interpolated. Left out one well at a time, no interpolator predicts a withheld well's residual appreciably better than the network mean, so no surface is drawn. The residual is interpreted as a structural diagnostic of where the model's three terms are sufficient, not as a quantified map of lateral inflow. This is because it also absorbs nonlinear soil-moisture storage dynamics, coefficient sensitivity and any systematic bias in the Thornthwaite PET estimate. The mechanistic attribution of the residual pattern is tested in two ways. An independent cross-correlation lag analysis tests for distance-dependent transport from the ridge. A seasonal climatology analysis tests for unmodelled summer evaporative demand.

3.11 Computational implementation

All data pre-processing, statistical modelling and spatial mapping were implemented as a sequential, numbered Python pipeline. Time-series management used pandas and NumPy, and model fitting used statsmodels. Spatial interpolation and rendering were carried out over a

LiDAR DEM. The state-space alignment and displacement calculation are implemented once, in a shared model-utilities module. Every script that fits the SSM calls this module, which ensures that all downstream analyses are consistent with the cluster-level fits. The pipeline and its outputs are reproducible by running the published scripts in sequence.

4. Results

4.1 Climate and well-network characterization

The RAF Valley record establishes that Newborough Warren operates under a persistent annual water surplus. Over the monitoring period, mean annual precipitation was approximately 891 mm and mean annual Thornthwaite PET approximately 652 mm, a study-period P/PET ratio of about 1.36 (Figure 2). The surplus is strongly seasonal. Winter rainfall exceeds evaporative demand by a wide margin. In summer, PET exceeds rainfall, so atmospheric demand on the aquifer is concentrated in the period of lowest recharge. Year-to-year variability in the water balance is driven almost entirely by precipitation, not by evaporative demand. Evaporative demand shows strong seasonality and a small positive inter-annual trend (1.96 mm/yr over 2007–2025), but it varies by only about 3% about its mean, against about 19% for winter rainfall. The post-2013 period has run warmer than the earlier summer-temperature baseline (Figure 3), which intensifies the summer deficit.

The reference network comprises 66 wells with monthly water-table records over the 2005–2026 monitoring period. It is supplemented by 22 extended-network wells with shorter records, classified by pattern matching (Section 4.3). The median record length is well above a decade. This gives the spatial density needed to distinguish cluster-scale behaviour from site-wide climate forcing.

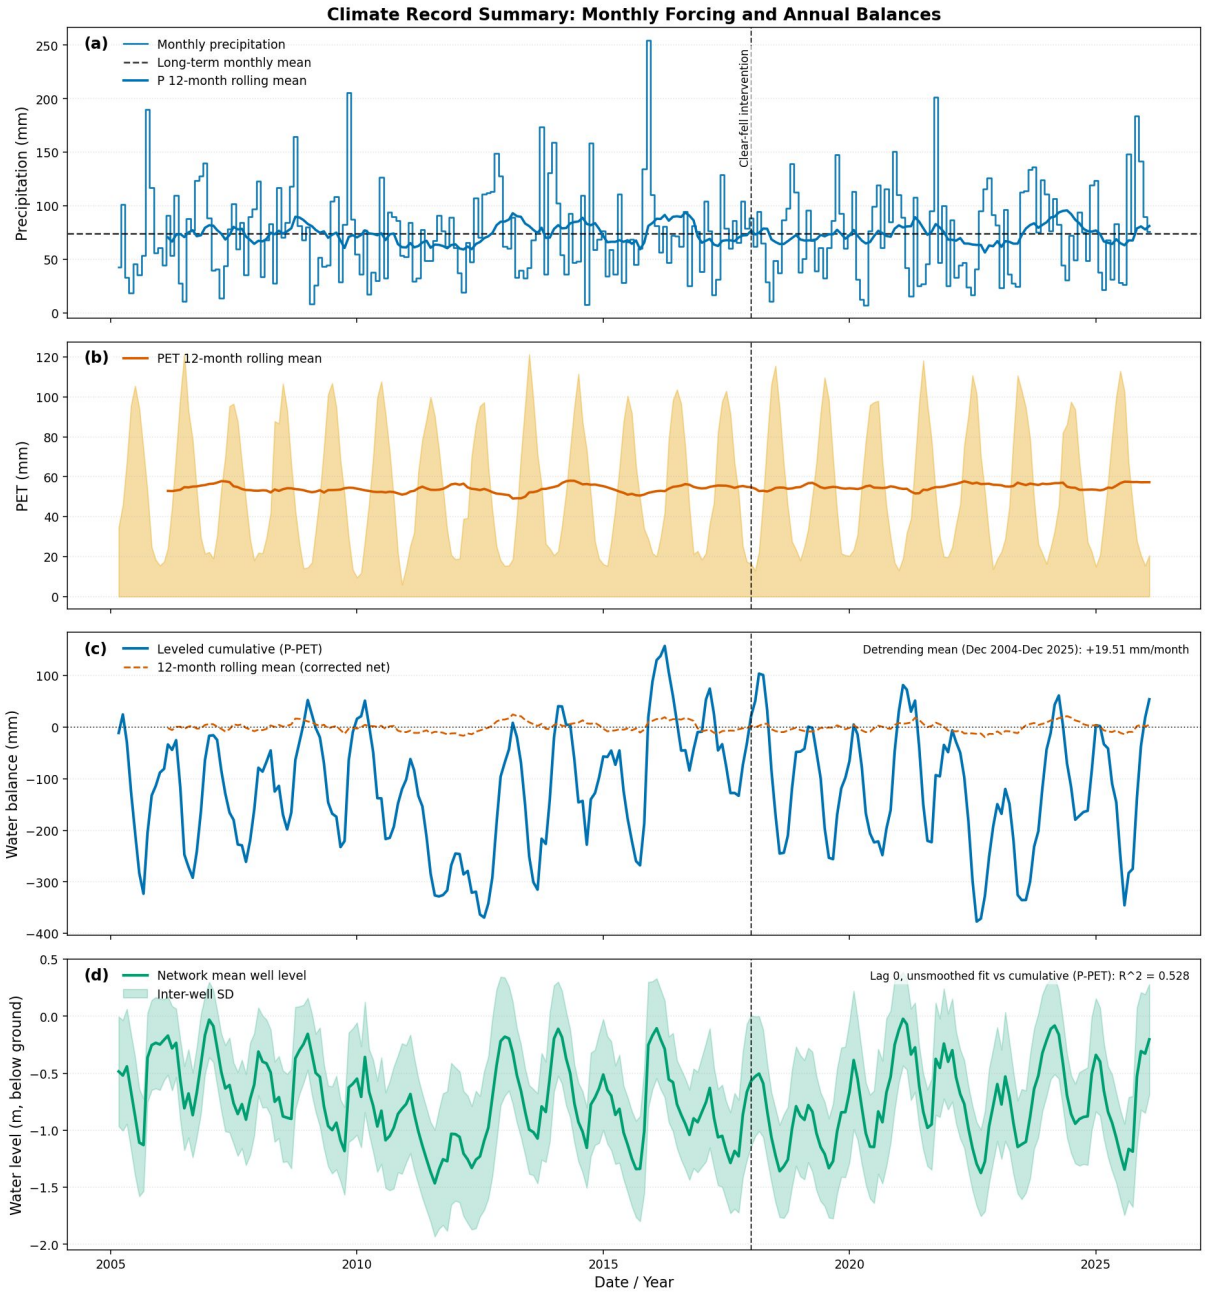


Figure 2. *Climate forcing over the monitoring period (2005–2026) from RAF Valley: monthly precipitation and Thornthwaite PET, with cumulative water surplus. The full ~95-year station record is not shown here; the long-term warming context is given in Figure 3.*

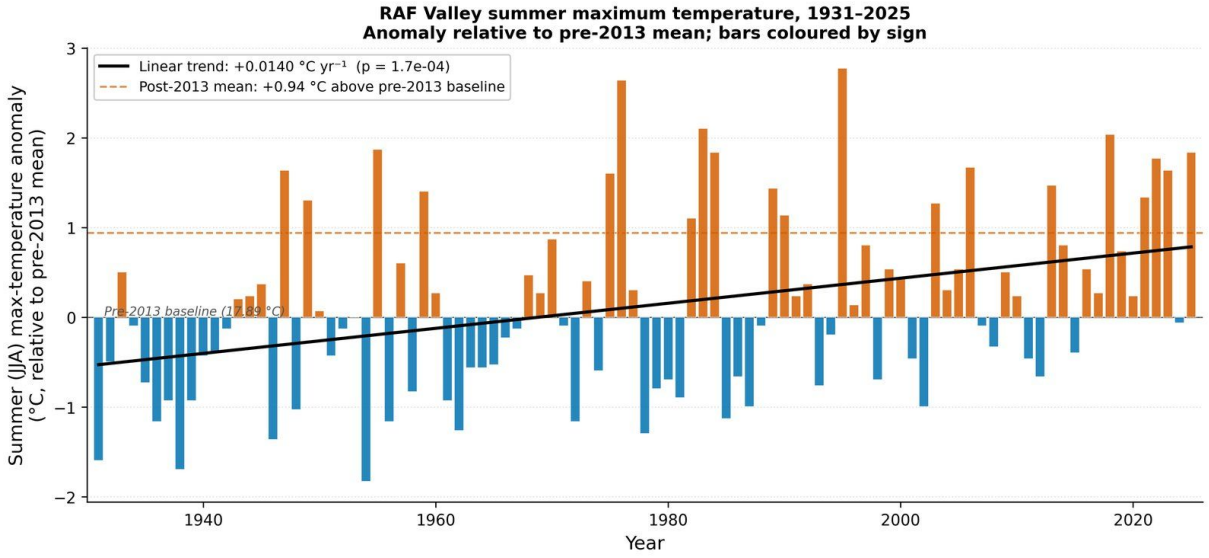


Figure 3. *Summer-temperature record showing the post-2013 warming relative to the earlier baseline.*

4.2 Network clustering and spatial architecture

Hierarchical clustering resolves five spatially coherent units under the dual-metric selection of k (Figure 4). The clustering used the standardized hydrograph dynamics, the Pearson correlation distance and Ward's linkage. The partition separates the network into a lake-margin unit (C1, Lake Edge), an open-dune unit (C2, Dune) and a western unit on deeper sand (C3, Western Residual). Two further units lie beneath the Corsican pine plantation and are distinguished by topographic position: C4, Main Forest, on the elevated ridge flank, and C5, Coastal Forest, on the lower coastal fringe (Figure 5). The cluster sizes are C1 = 7, C2 = 19, C3 = 26, C4 = 9 and C5 = 5 wells.

The five units recover the broad substrate contrast between the responsive, storage-limited east and the buffered, deeper-sand south-west (Stratford et al., 2007). The south-western unit (C3) is gradational toward both the open dune and the forest, rather than a sharply bounded block. The top Ward merge (linkage distance ≈0.99) separates C4 Main Forest from the other 57 wells. The east–west division, C1 and C2 against C3 and C5, follows at ≈0.61. C5 Coastal Forest therefore joins the western group rather than C4, despite sharing the pine canopy. Figure 6 shows the cluster-centroid hydrographs and seasonal amplitudes.

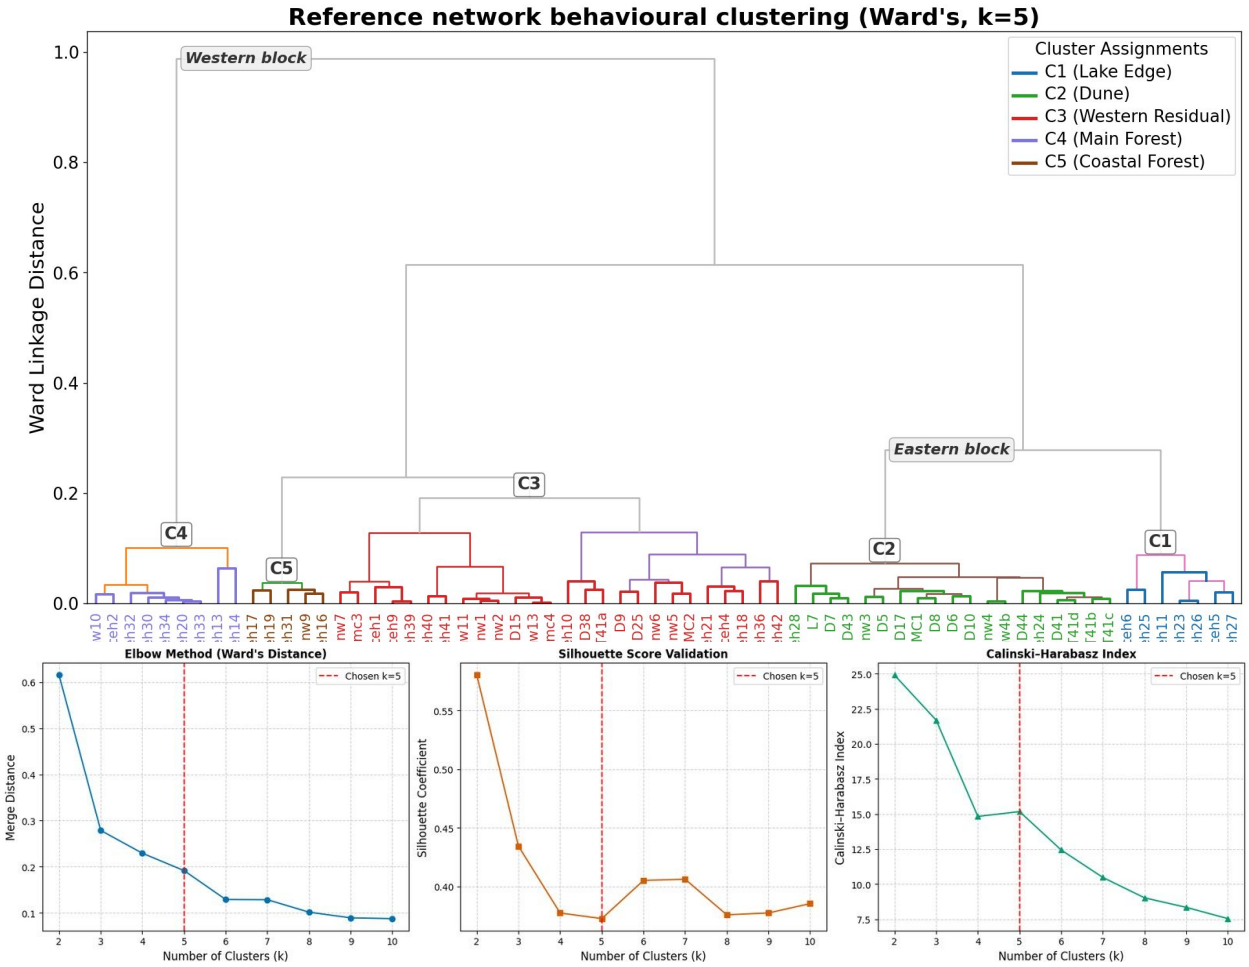


Figure 4. *Hierarchical clustering of hydrograph dynamics: Ward's dendrogram (upper) and the dual-metric cluster-number selection (lower).*

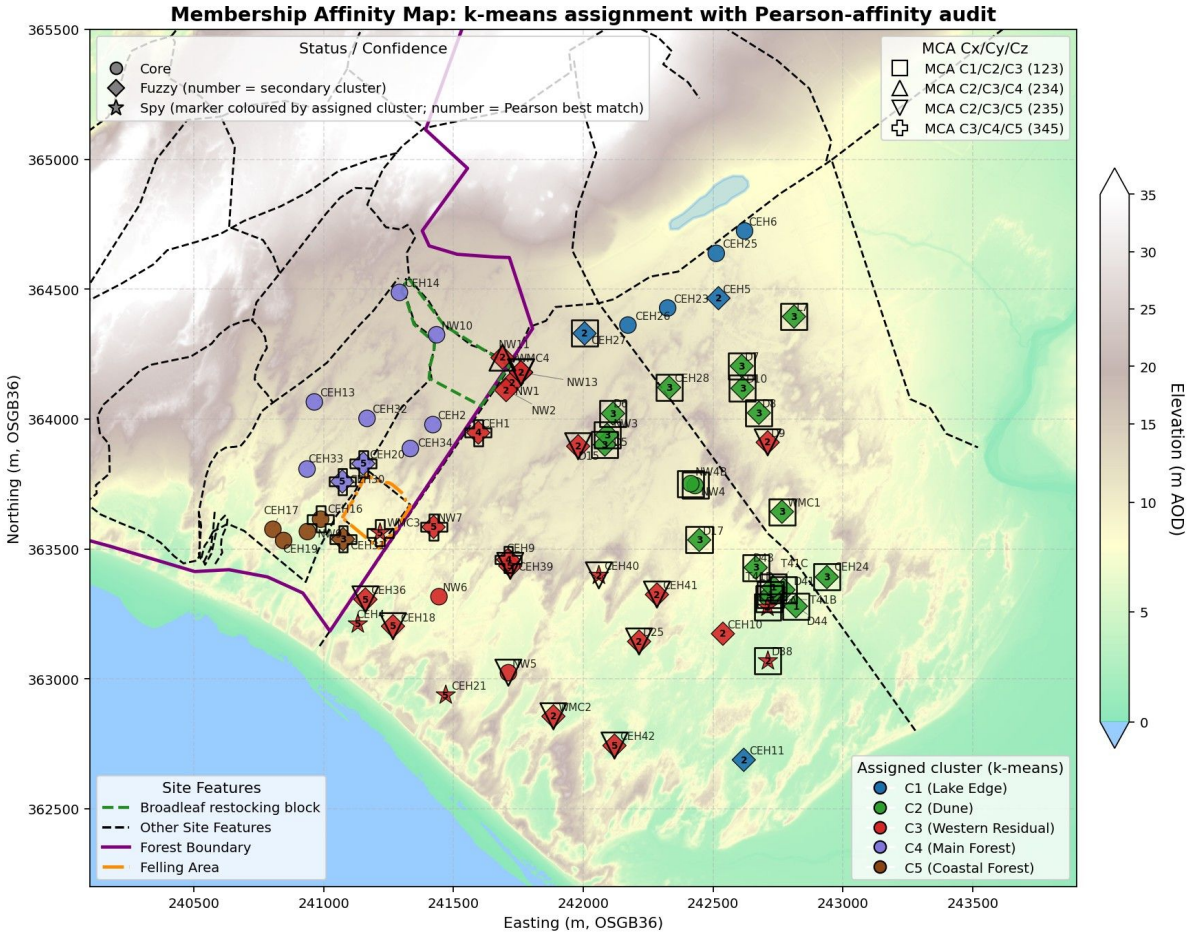


Figure 5. *Spatial distribution of the five hydrogeological clusters (Pearson affinity validation), reference network (66 wells). Points coloured by k-means assignment; circles = Core (Δr > 0.05), diamonds = Fuzzy (Δr < 0.05; number inside diamond = secondary cluster), stars = Spy (Pearson best-match ⁄ assigned; number inside star = best-match cluster). Hollow overlays = Multi-Cluster Affinity (r > 0.90 with ≥3 clusters). LiDAR: © NRW & OS.*

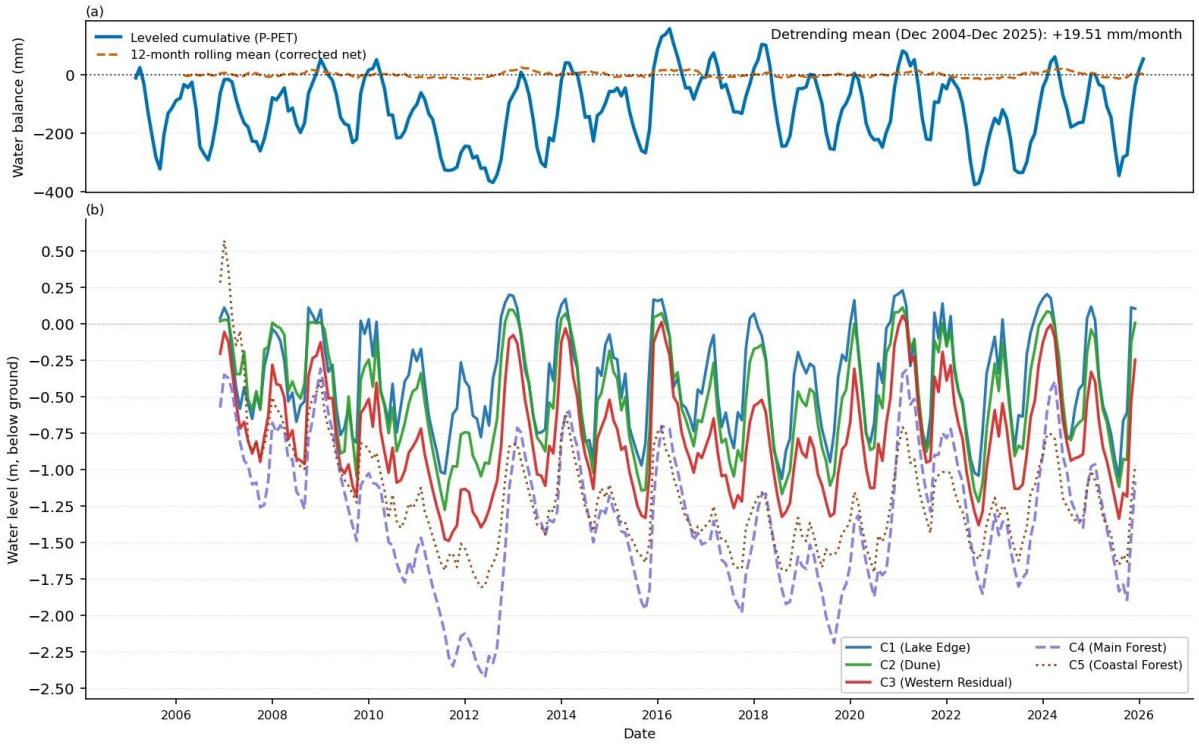


Figure 6. *Cluster-centroid hydrographs and seasonal amplitudes, 2005–2026.*

4.3 Affinity validation of the cluster structure

The Pearson affinity analysis supports the primary membership of 60 of the 66 reference wells. Membership falls into tiers, set by the correlation delta between the assigned cluster and the next-strongest one. A minority of wells are core members, with an unambiguous dominant affinity. Many others show closely competing affinities between neighbouring units. This is consistent with gradational rather than sharp boundaries in a contiguous, highly transmissive sand aquifer (Rao and Srinivas, 2006). Multi-cluster-affinity flags are concentrated along the C1/C2/C3 transition and the forest–dune margin. This concentration identifies the zones where geological and vegetative controls overlap.

4.4 Cluster mechanistic coefficients

The displacement-formulation SSM, fitted to each cluster-centroid hydrograph, characterizes the five clusters mechanistically (Table 1). All coefficients are statistically significant ($p < 0.003$) at every cluster.

Table 1. *Cluster mechanistic characterization. β_1 = recharge sensitivity (mm rise per mm rainfall); β_2 = atmospheric draw (mm decline per mm PET); $-\beta_3$ = proportional drainage decay (fraction of displacement above the 3.7 m datum dissipated per month); LCSC = $100/\beta_1$, the reciprocal of β_1 (mm of rainfall per 100 mm of water-table rise); R^2 = one-step diagnostic fit.*

Cluster	n wells	β_1	β_2	$-\beta_3$	LCSC (%)	R^2
C1 Lake Edge	7	4.578	0.911	0.089	21.844	0.734

C2 Dune	19	3.896	1.666	0.063	25.669	0.745
C3 Western Residual	26	3.758	1.778	0.062	26.611	0.825
C4 Main Forest	9	2.476	2.570	0.018	40.392	0.733
C5 Coastal Forest	5	2.412	1.248	0.045	41.462	0.677

Recharge sensitivity β_1 is highest at the Lake Edge (C1: 4.58). This indicates that each millimetre of rainfall drives a proportionally larger water-table rise in the shallow-till eastern aquifer than in the Dune (C2: 3.90) or the Western Residual (C3: 3.76). The two forested clusters return near-identical and substantially lower recharge sensitivities (C4: 2.48, C5: 2.41). This is consistent with canopy interception suppressing the rainfall signal that reaches the water table. Expressed as LCSC ($100/\beta_1$), about 22–26 mm of rainfall raises the water table by 100 mm at the Lake Edge and Dune. The figure is 27 mm in the Western Residual and 40–41 mm beneath the plantation. Beneath the canopy, nearly twice as much rain is therefore needed for the same rise.

Atmospheric draw β_2 is highest at C4 (2.56). There it slightly exceeds C4's own β_1 (2.56 against 2.48), so the Main Forest cluster loses marginally more water table per millimetre of PET than it gains per millimetre of rainfall. C3 (1.78) and C2 (1.67) follow, and C5 (1.25) and C1 (0.91) are lowest. C5 carries the same pine canopy as C4 yet returns less than half its atmospheric draw. This points to topographic position and water-table depth, rather than canopy, as the dominant control of the cluster-level β_2 contrast (Section 4.8). C1's low β_2 reflects its shallow water table and rapid lateral drainage to the lake, where head-dependent drainage captures the dominant loss.

The drainage decay $-\beta_3$ shows the strongest contrast of the three coefficients. It is highest at C1 (0.089), where rapid lateral exchange with Llyn Rhos-Ddu drives fast recession. It is similar at C2 (0.063) and C3 (0.062), and lower at C5 (0.045). It is lowest at C4 (0.018), where a suppressed recharge pulse and weak lateral discharge produce the most sluggish drainage in the network.

4.5 Water balance and specific yield

The head-space water-balance decomposition (Table 2; Figure 7) closes to within 1.8% of total losses at every cluster. This closure is a property of the fit, not independent evidence. A model fitted to heads alone can be expected to close the mean balance (Hollingham, 2026c). The closure therefore shows that the three terms are mutually consistent at the cluster mean without an unmeasured boundary term. It does not show that each flux is confirmed. Recharge declines systematically from C1 (0.340 m/month) to the forested clusters (≈0.18 m/month). Total losses track recharge closely at every cluster.

The drainage partition varies markedly. Drainage accounts for 85% of total losses at C1 but only 24% at C4, where atmospheric draw dominates. That split is conditional on the datum and is not identified from heads alone. Consider the range from the shallowest datum at which every cluster's drainage term is identified to a datum of 8 m. Across it, the drainage share runs from 73% to 89% at C1 and from 4% to 28% at C4, with the clusters in the same order throughout (Hollingham, 2026c).

Table 2. *Mean monthly head-space water balance (m/month), 2005–2026. Identical climate forcing at every cluster ($\bar{P}$ ≈ 74.4 mm/month, $\bar{PET}$ ≈ 54.5 mm/month). The drainage/ET split is taken at the 3.7 m datum and is conditional on it (Section 4.5).*

Cluster	Recharge	Atm. draw	Drainage	Total loss	Residual	Drain % / ET %
C1 Lake Edge	0.341	0.050	0.295	0.345	−0.004	86 / 14
C2 Dune	0.290	0.090	0.203	0.293	−0.003	69 / 31
C3 Western Residual	0.278	0.097	0.181	0.278	<0.001	65 / 35
C4 Main Forest	0.184	0.140	0.044	0.184	<0.001	24 / 76
C5 Coastal Forest	0.179	0.068	0.115	0.183	−0.003	63 / 37

The indicative volumetric conversion (Table 3) places all clusters under ≈892 mm/yr rainfall. At the forested clusters, canopy interception removes ≈214 mm/yr, leaving ≈678 mm/yr of throughfall. The mid-point of the ET/drainage partition ranges from drainage-dominated at C1 (≈651 mm/yr drainage vs ≈242 mm/yr ET) to ET-dominated at C4 (≈252 vs ≈426 mm/yr). The ET column counts only what evaporates after the canopy has taken its share. Interception is itself evaporation, so a forest's total evaporative loss is I + ET: ≈641 mm/yr at C4, close to PET (≈655 mm/yr), against ≈440 mm/yr at C2. At C5 the total (≈428 mm/yr) is similar to C2's. C2's mid-point is the least certain in the table, because the two partitioning methods disagree most there (ET ≈275–606 mm/yr). At C2 and C4 the recession estimate lies outside the range the model gives across the admissible datum (1.6–8 m; the bar beside each loss bar in Figure 7b), so the split there depends on the method more than on the datum.

Table 3. *Indicative annual volumetric water balance (mm/yr). P is the RAF Valley rainfall, the same series for every cluster, averaged over the months in which that cluster has readings; because the network was installed progressively those months differ slightly, so P varies from 889 to 893 mm/yr. I = canopy interception (24% of P; Freeman, 2008), forested clusters only. ET is the evaporative loss after interception, from throughfall and the water table; because I is itself evaporation, the total evaporative loss is I + ET, computed before rounding, so it can differ by one from the sum of the rounded columns. ET/drainage are mid-points of two independent partitioning methods; the SSM method is taken at the 3.7 m datum and is conditional on it (Section 4.5).*

Cluster	P	I	P_net	ET (mid)	I + ET (mid)	Drainage (mid)
C1 Lake Edge	893	0	893	242	242	651
C2 Dune	892	0	892	440	440	452
C3 Western Residual	889	0	889	289	289	601
C4 Main Forest	893	214	679	426	641	252
C5 Coastal Forest	892	214	678	214	428	464

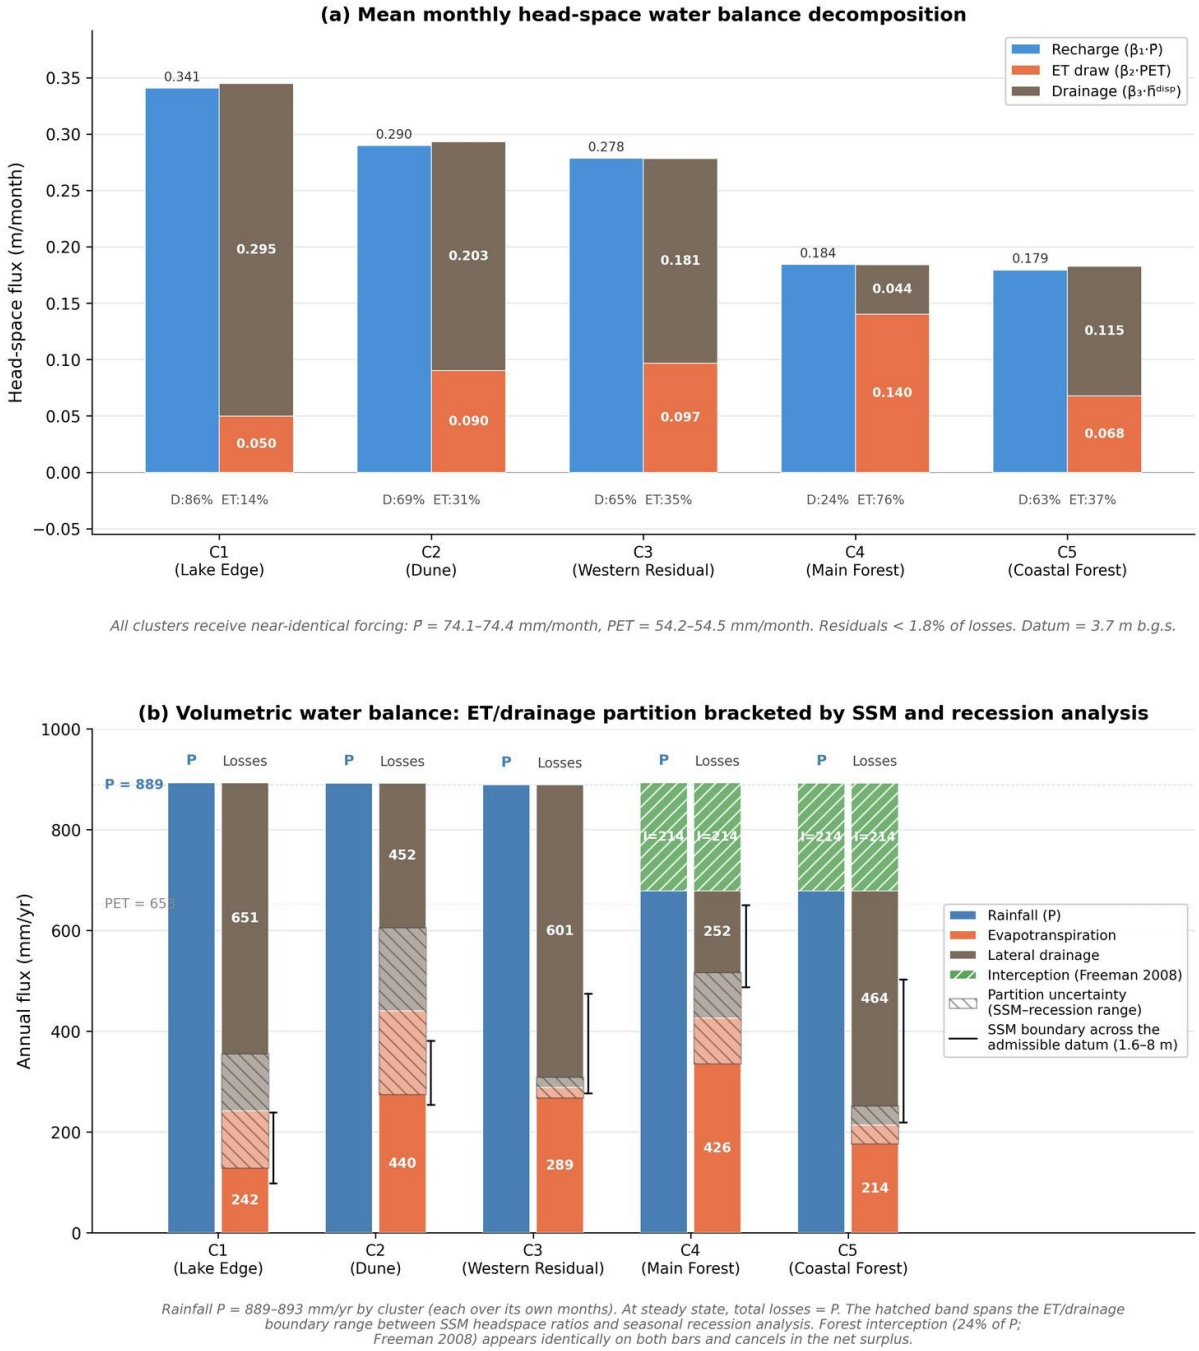


Figure 7. *Water-balance decomposition by cluster: head-space (recharge versus atmospheric draw and drainage) and indicative volumetric partition. In (b), the hatched partition-uncertainty band spans the two independent estimates of the ET/drainage split — the SSM head-space ratio and the seasonal winter/summer recession ratio; it is smallest for C3, where the two methods agree to within 5 percentage points (drainage fractions 0.70 and 0.65). The black bar beside each loss bar spans where the SSM ET/drainage boundary falls as the drainage datum is varied across its admissible range (1.6–8 m); it shows how far the SSM estimate depends on the datum and is not a statistical confidence interval.*

The water-table-fluctuation (WTF) specific-yield estimates (Table 4; Figure 8) provide an independent test. They ask whether the cluster contrasts reflect storage architecture or surface

boundary conditions. Uncorrected, the forested-cluster event-median specific yields sit at or above the top of the open-dune range (C4: 0.315, against an open-dune maximum of 0.325; C5: 0.358, above it). This is because gross rainfall overstates the recharge flux beneath the canopy. Applying the canopy-interception correction to the recharge term brings the cluster-level event medians down to 0.261 (C4) and 0.330 (C5). C4 then falls cleanly into the open-dune range (0.210–0.325), while C5 sits just above its upper edge. Aggregating instead as the median of the per-well estimates gives 0.253 and 0.308. The two aggregations are reported separately throughout and are not interchangeable.

The correction uses a single literature fraction rather than a fitted one, so the convergence is tested against that fraction directly. The forest clusters' event medians were re-evaluated over interception fractions from 0 to 0.40 (Section 3.8). Both fall within the open-dune range only for fractions between 0.26 and 0.34: C4 for every fraction up to 0.34, C5 only from 0.26. The C4 result therefore does not depend on the particular value 0.24, while C5 sits just outside the range at that value. The seasonal residual analysis of Section 5.4 is an independent check, and its cross-cluster ordering is likewise inconsistent with a mis-set fraction.

Across the three open-dune clusters, the per-well specific yields are well described by a planar spatial trend ($n = 52$, $R^2 = 0.610$, residual standard deviation 0.031). The trend rises by 0.070 per kilometre along an azimuth of 231°, so storage increases steadily toward the south-west. Both components are significant (easting −0.054 per kilometre, northing −0.044 per kilometre, each $p < 0.001$). The trend is fitted to, and reported for, the open-dune network only. The forested wells lie outside its spatial domain, and the surface is not extrapolated beneath the plantation.

The C4 result indicates that the attenuated water-table response beneath the plantation reflects the surface canopy boundary condition rather than a fundamentally different substrate. The C5 estimate is the least reliable in the network. A majority of its rising-limb events imply an event-level Sy at or above the porosity of dune sand. Its median is therefore constrained by the physical-plausibility limit rather than freely estimated, and it is reported as a weak corroborator only. All WTF medians exceed the assumed values used in the volumetric balance by a factor of two to three. They are interpreted as upper bounds (Healy and Cook, 2002; Scanlon et al., 2002).

Table 4. *Specific yield by cluster, water-table-fluctuation method. Event-median Sy; forested clusters reported uncorrected and interception-corrected.*

Cluster	Sy assumed	Sy WTF median (uncorr.)	Sy WTF median (corrected)
C1 Lake Edge	0.08	0.210	—
C2 Dune	0.12	0.255	—
C3 Western Residual	0.12	0.325	—
C4 Main Forest	0.12	0.315	0.261
C5 Coastal Forest	0.12	0.358	0.330

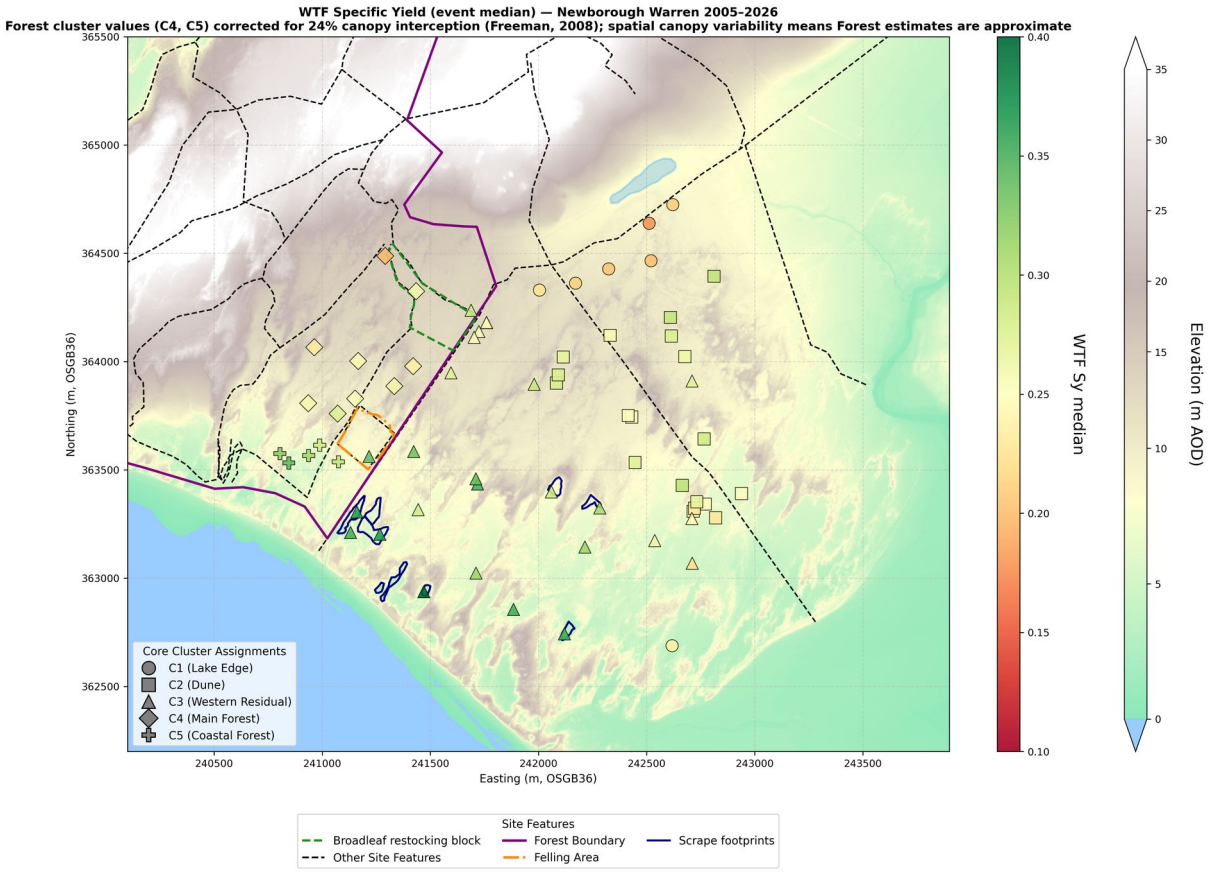


Figure 8. *Specific yield at each reference well, water-table-fluctuation method (the median of the well's events); forest wells (C4, C5) interception-corrected. Marker shape gives the cluster and colour the specific yield, over LiDAR elevation. Values change gradually across the network: neighbouring wells either side of a cluster margin differ about as much as neighbours within a cluster, so there is no step at the margins. The values are upper bounds from a lumped-parameter model, not calibrated aquifer properties.*

4.6 Model benchmarking: SSM versus traditional linear model

Run freely over the comparison window, the SSM raises the median iterative Nash–Sutcliffe efficiency from 0.42 under the TLM to 0.77. It returns a positive efficiency at 65 of the 66 reference wells, against 55 for the TLM (Hollingham, 2026c). The gain is not uniform across the network (Figure 9). It is largest in the lake-proximal C1 Lake Edge, where head-dependent drainage is strongest (cluster median +0.45, against +0.38 at C2 and +0.37 at C3; e.g. CEH6: TLM −0.35 → SSM 0.71). It is smallest in the C4 forest interior (+0.12), where drainage is weakest. At two wells there, the SSM scores below the TLM. The spatial pattern of the gain therefore serves as a non-invasive diagnostic of where head-dependent drainage governs the aquifer.

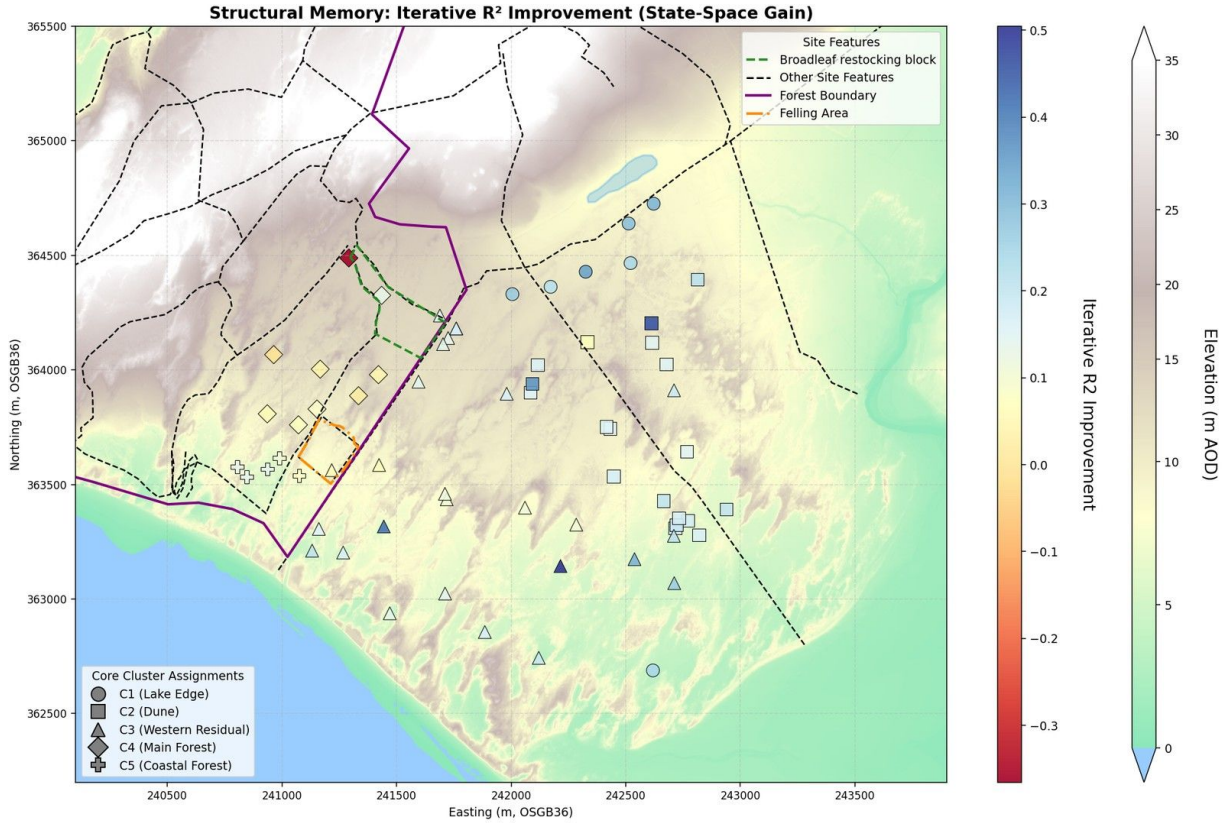


Figure 9. *Spatial pattern of the SSM-over-TLM iterative Nash–Sutcliffe improvement across the reference network; the models and the free-running protocol are set out in Hollingham (2026c).*

4.7 Spatial coefficient structure and drainage half-life

Optimized well by well, the R^2-maximizing datum tracks each well's mean water-table depth (r = 0.59). It follows the ground rather than a fixed elevation. Regressing its elevation on ground elevation across the 66 reference wells gives a slope of +0.845; a base at constant depth below ground would give +1, and one at fixed elevation 0. The R^2 gain over the uniform 3.7 m datum is small across most of the network (median 0.014; Figure 10). It exceeds 0.05 at only eight wells, at the Lake Edge and in the eastern dune, where the water table is shallowest. Because the fit-optimal datum is confounded with mean water-table depth, it does not locate the impermeable base. What the sweep can and cannot decide is examined in Hollingham (2026c).

Interpolating the per-well coefficients yields continuous surfaces (per-well ranges in Table 5; Figures 11–13). Their broad structure aligns with the cluster boundaries, and they preserve substantial within-cluster variation. Recharge sensitivity β_1 is highest in the north-east around the Lake Edge wells (per-well up to 6.24). It declines across the open dune to its lowest values beneath the forest. Atmospheric draw β_2 peaks in the C4 forest interior and along the ridge-adjacent northern dune, and is lowest at the Lake Edge. Drainage rate β_3 shows the strongest spatial contrast of any coefficient. It runs from the fast-draining Lake Edge (per-well up to 0.134

per month) to the near-stagnant forest interior. There, one ridge-flank well (CEH14) records a negative value, where the displacement formulation does not capture the boundary condition.

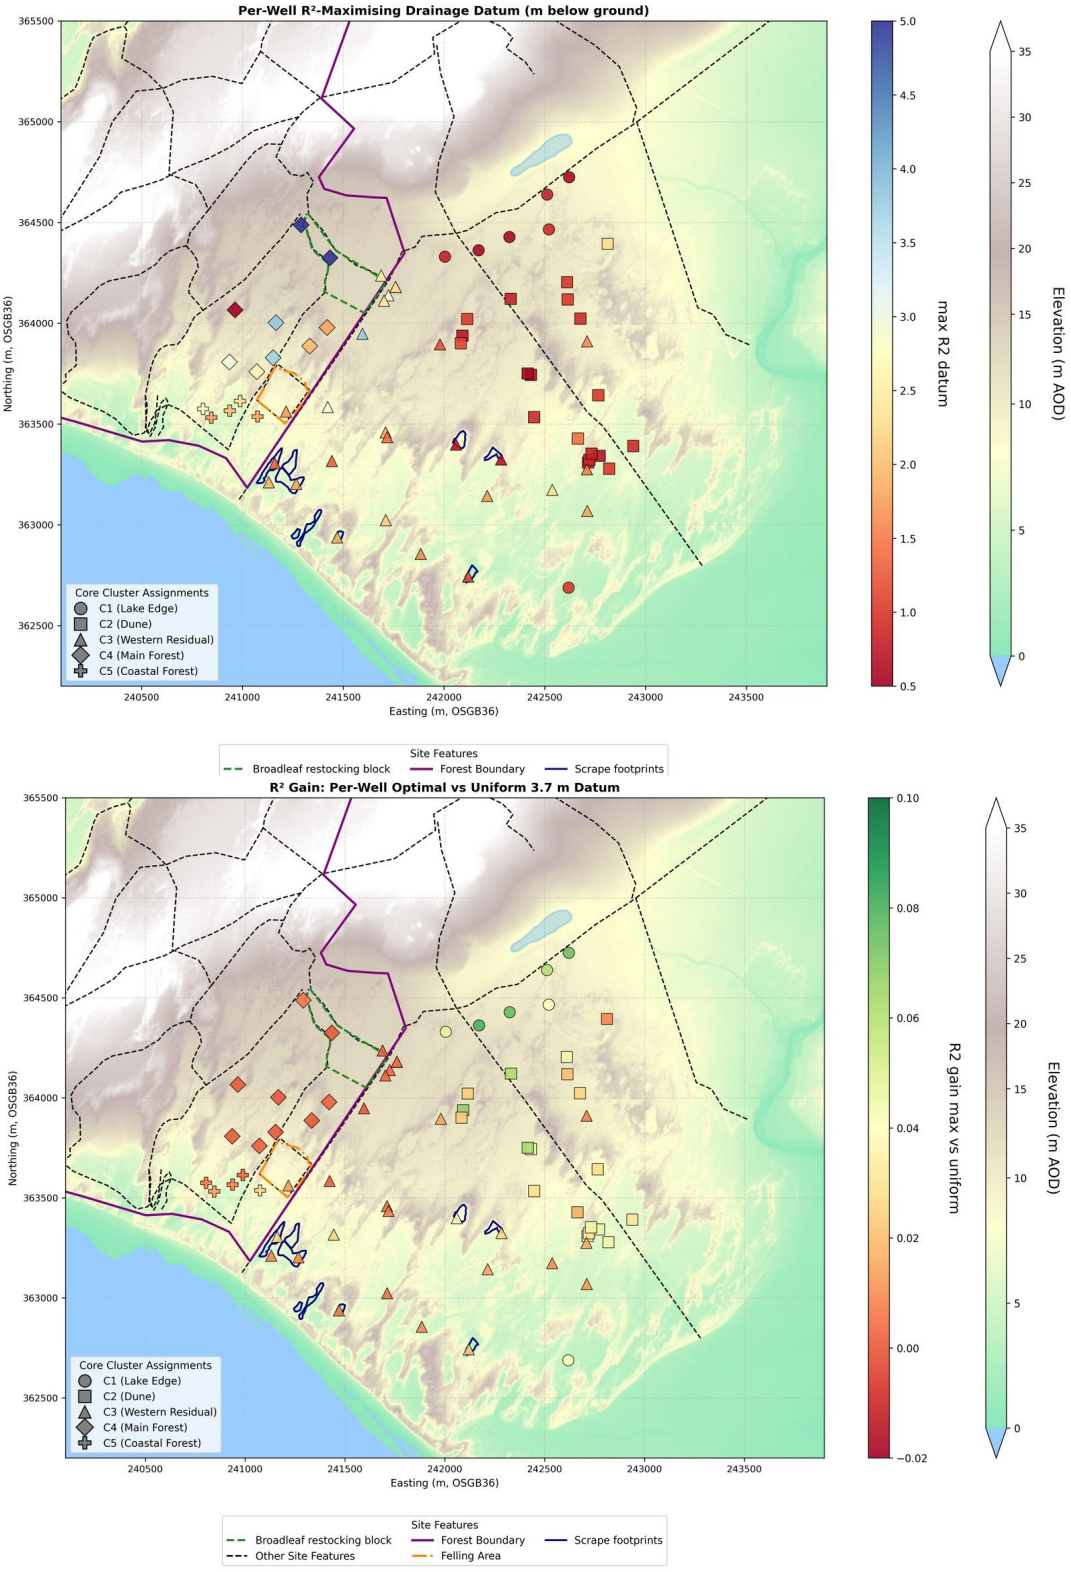


Figure 10. *Per-well optimal drainage datum (top) and the R^2 gain over the uniform 3.7 m datum (bottom).*

Table 5. *Per-well SSM coefficient ranges by cluster.*

Cluster	β_1 range	β_2 range	β_3 range
C1 Lake Edge	4.267–6.237	−0.400–1.251	0.088–0.134
C2 Dune	3.439–4.716	1.034–2.531	0.057–0.090
C3 Western Residual	2.332–5.514	0.884–2.185	0.034–0.109
C4 Main Forest	2.072–3.481	2.062–3.792	−0.020–0.040
C5 Coastal Forest	2.125–2.752	0.776–1.359	0.039–0.059

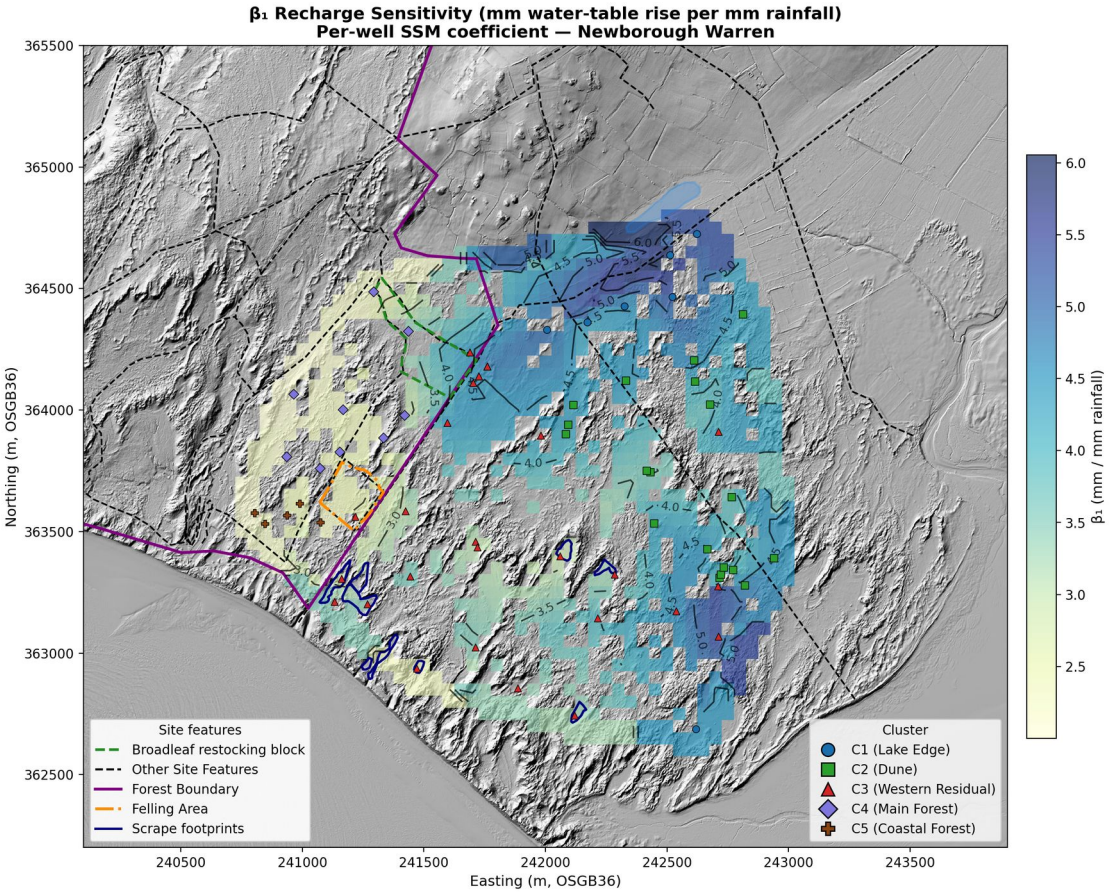


Figure 11. *Spatial SSM coefficient atlas for the 66-well reference network, Newborough Warren.* ***β_1 recharge sensitivity*** *(mm water-table rise per mm rainfall) — highest at the Lake Edge (C1, 4–6 mm/mm), with a sharp drop at the forest boundary consistent with 24% canopy interception. Per-well coefficients interpolated by piecewise-linear (Delaunay) interpolation over 2 m LiDAR DEM hillshade, with ridge and non-aquifer areas masked. Well markers coloured by cluster assignment. (LiDAR data: © NRW & OS)*

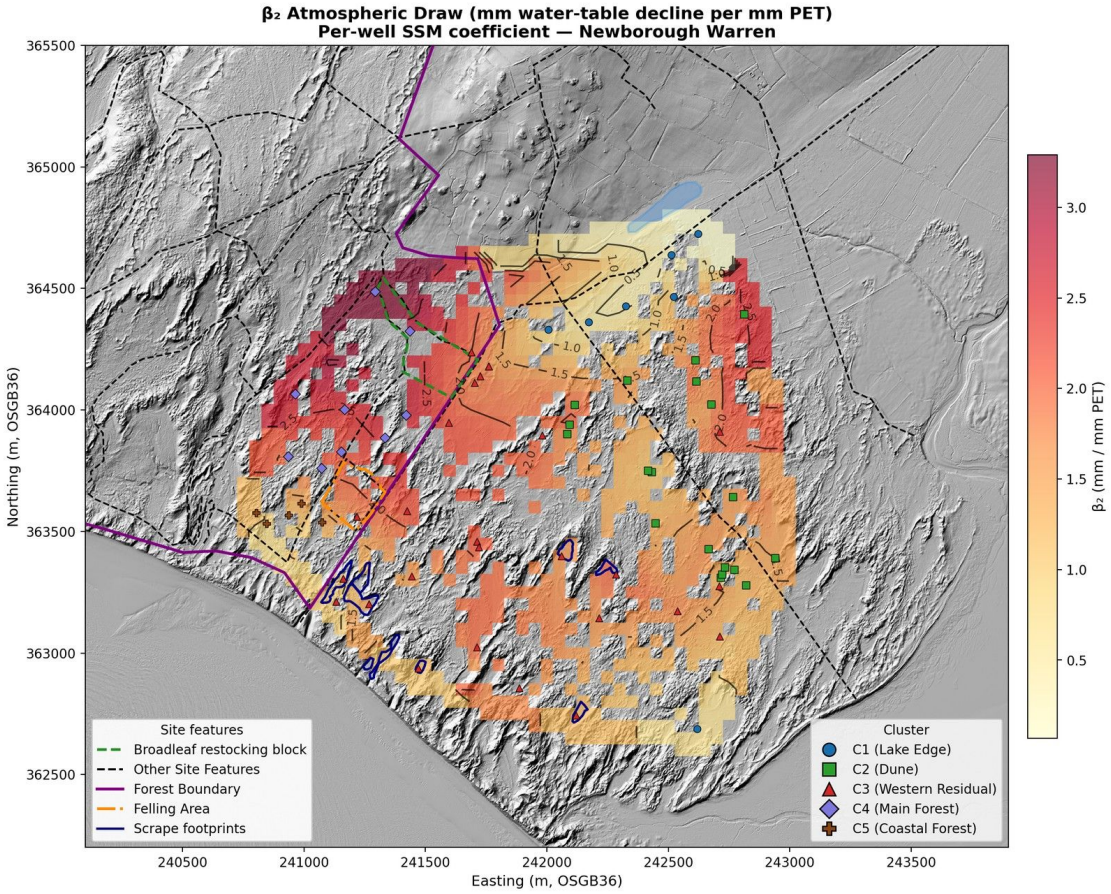


Figure 12. *Spatial SSM coefficient atlas for the 66-well reference network, Newborough Warren.* ***β₂ atmospheric draw*** *(mm water-table decline per mm PET) — highest in the forest interior (C4, 2.1–3.8 mm/mm), lowest at the Lake Edge (C1) and Coastal Forest (C5). Per-well coefficients interpolated by piecewise-linear (Delaunay) interpolation over 2 m LiDAR DEM hillshade, with ridge and non-aquifer areas masked. Well markers coloured by cluster assignment. (LiDAR data: © NRW & OS)*

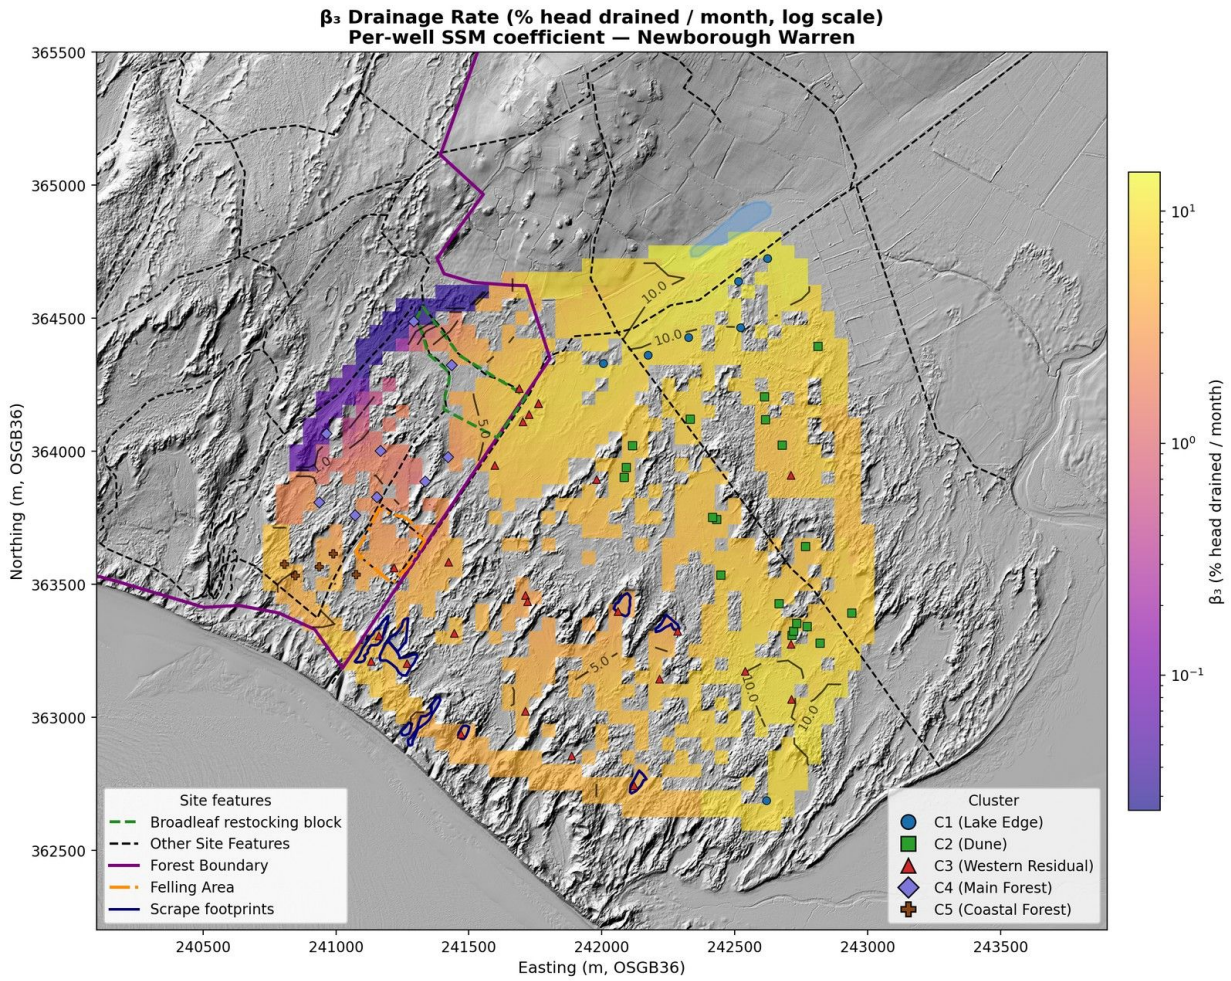


Figure 13. *Spatial SSM coefficient atlas for the 66-well reference network, Newborough Warren.* ***β_3 drainage rate*** *(% displacement drained per month, logarithmic scale) — a sevenfold contrast in cluster-mean drainage rate, from the fast-draining Lake Edge (C1, 9–13% per month) to the near-stagnant forest interior (C4, no well above 4.0% per month). CEH14, the sole well with negative β_3, sits at the ridge flank. Per-well coefficients interpolated by piecewise-linear (Delaunay) interpolation over 2 m LiDAR DEM hillshade, with ridge and non-aquifer areas masked. Well markers coloured by cluster assignment. (LiDAR data: © NRW & OS)*

The drainage decay half-life $t_{½} = \ln(2)/\beta_3$ is the time that modelled storage above the datum takes to halve under the drainage term alone (Table 6; Figure 14). It is the recession constant of the drainage term at the project datum, and it scales with the datum (Table 6). It is not the time the water table takes to return to its mean after a disturbance. That time is measured by the free-intercept form, and it is shorter by more than a factor of two at the median well (Hollingham, 2026c). The half-life is a timescale, not a bound on amplitude. Over a summer, drainage removes only a small part of a winter excess, even where the half-life is long. The much larger seasonal fall is set by evapotranspiration (the β_2 term).

The eastern open dune drains fastest (C1 median 8.0 months, C2 10.7). The western dune is intermediate (C3 median 11.6 months). The forest interior drains most slowly in the network (C4 median 23.0 months, individual wells to 42). C5 Coastal Forest (median 14.1 months) resembles the western dune despite sharing the pine canopy, which reinforces the substrate

and elevation control. The half-life is set by the drainage rate alone. The storage-drainage index $\tau = Sy/\beta_3$ additionally weights it by specific yield. Sy and $1/\beta_3$ are only weakly correlated across the network ($r = 0.22$, $p = 0.09$). The two therefore produce closely similar spatial patterns ($r = 0.94$), and the cluster ranking is unchanged.

Table 6. *Drainage decay half-life $t_{½} = \ln(2)/\beta_3$ (months) by cluster. Ridge-flank wells CEH14 (negative β_3) and CEH13 (near-zero β_3) excluded. β_3 is the recession constant of the drainage term at the 3.7 m displacement datum, so $t_{½}$ is conditional on that reference rather than a datum-free property of the aquifer. Fitting the same model on recession months alone reproduces every cluster's coefficient (ratios 1.01 to 1.24, per-well $r = +0.944$). Over ±0.5 m about the datum (3.2–4.2 m; the lower bound is set by the deepest recorded water table, 2.87 m at CEH32 in June 2012, below which the displacement changes sign) every cluster half-life rises with the reference. The cluster ranking is invariant across that range and far beyond it (C1 < C2 < C3 < C5 < C4 at every datum from 2.9 to 8.0 m). The ratios are not: relative to C3 they drift by 0.2% (C2), 2.4% (C1), 6.1% (C5) and 10.5% (C4) over 3.2–4.2 m. Absolute values and inter-cluster ratios are therefore conditional on the project datum; the ordering is not.*

Cluster	t½ min	t½ median	t½ max	n
C1 Lake Edge	5.9	8.0	8.9	7
C2 Dune	8.0	10.7	12.8	19
C3 Western Residual	6.1	11.6	19.4	26
C4 Main Forest	16.3	23.0	41.6	7
C5 Coastal Forest	11.3	14.1	15.7	5

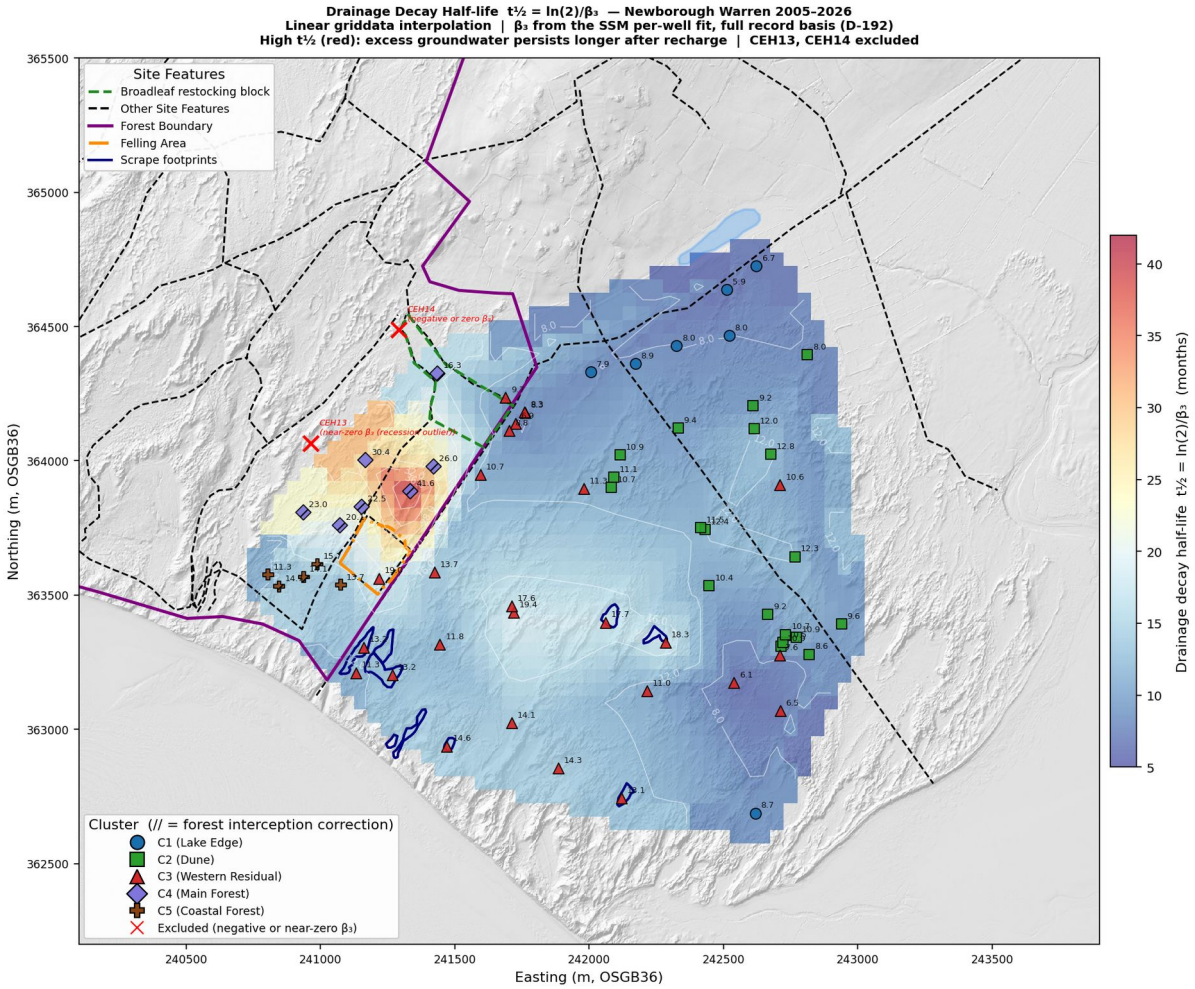


Figure 14. *Drainage decay half-life t½ = ln(2)/β₃ across the reference network. The interpolation extends over the rock-ridge bedrock outcrop on the northern boundary, which carries no monitoring wells; values shown there are extrapolations from the surrounding network and are not interpretable as model output on bedrock.*

4.8 Substrate control within the forested zone

Within the forested zone the per-well coefficients span much of the site-wide range. Regression against elevation, distance from the ridge and easting identifies the underlying control (Table 7). Atmospheric draw β_2 is overwhelmingly governed by elevation ($r = 0.989$, $p < 0.001$; $R^2 = 0.978$): higher wells on the ridge return systematically higher β_2. Drainage β_3 shows a weaker but significant elevation dependence ($r = -0.831$, $p < 0.001$), with distance from the ridge adding information (combined $R^2 = 0.783$). Recharge β_1 shows no significant elevation relationship ($p = 0.57$); easting is its best single predictor ($r = 0.56$, $p = 0.039$), and its correlation with distance from the ridge ($r = -0.49$) is not significant ($p = 0.076$).

β_2 does vary substantially across clusters (Table 1), but the cluster-level contrast is set by substrate, water-table depth and topographic position rather than by canopy maturation: the C4–C5 split, a 2-fold difference in atmospheric draw within the plantation, is the cleanest evidence. C4 and C5 form two distinct groups rather than a continuum. They separate almost entirely on β_2 — C4 at 2.06–3.79, C5 at 0.78–1.36, with no overlap — and on mean elevation

(C4: 10.6 m, C5: 5.6 m; p = 0.0001), corresponding to the transition from the elevated ridge flank, with its deeper water table and unsaturated zone, to the low coastal fringe, with a shallower water table. The canopy is not uniform. Its density was measured at each forest well from the dated aerial imagery, as a canopy texture index (0 for open ground, 1 for mature conifer) in a 30 m disc around the well, on the frames before the 2017 felling. Across the fourteen wells the index ranges from 0.15 to 1.47, yet it is not significantly related to β_2 (r = 0.26, p = 0.37) or to elevation (r = 0.25, p = 0.39), while β_2 follows elevation (r = 0.989). Holding elevation constant, canopy adds nothing (partial r = 0.09, p = 0.77), and the result is the same with 20 m and 50 m discs. C4 and C5 do not differ in canopy (medians 0.99 and 0.82). The C4–C5 separation therefore tracks elevation and water-table depth rather than canopy. The subsurface beneath the plantation is not cored; on the topographic and coefficient evidence the C4/C5 split is best read as a substrate-and-elevation discontinuity rather than an arbitrary division of a continuous canopy population.

Table 7. *Within-forest spatial predictors of per-well SSM coefficients (n = 14). Pearson r (p).*

Coefficient	vs Elevation	vs Dist. from ridge	vs Easting
β_1 recharge	0.167 (0.568)	−0.489 (0.076)	0.556 (0.039)
β_2 atm. draw	0.989 (<0.001)	−0.909 (<0.001)	0.744 (0.002)
β_3 drainage	−0.833 (<0.001)	0.651 (0.012)	−0.484 (0.080)

4.9 Mean water-table surface and flow field

The water table is a single mound, highest beside the bedrock ridge (mean head 13.32 m AOD at CEH14). It falls to the coastal head along Caernarfon Bay and to the river and lake levels inland (Figure 15). Left out one at a time, the wells are predicted with a median absolute error of 0.14 m in the mean state and almost no bias (−0.005 m). Flow runs outward from the high point. It runs south and south-west to Caernarfon Bay across the west and centre, north towards Llyn Rhos-Ddu from the slacks beside the lake, and east and south-east to the Afon Braint. The pattern barely changes with the state of the aquifer. The median gradient is 0.0071 in the wet state and 0.0062 in the dry, and the median flow direction at the slacks turns 4.7° between the two.

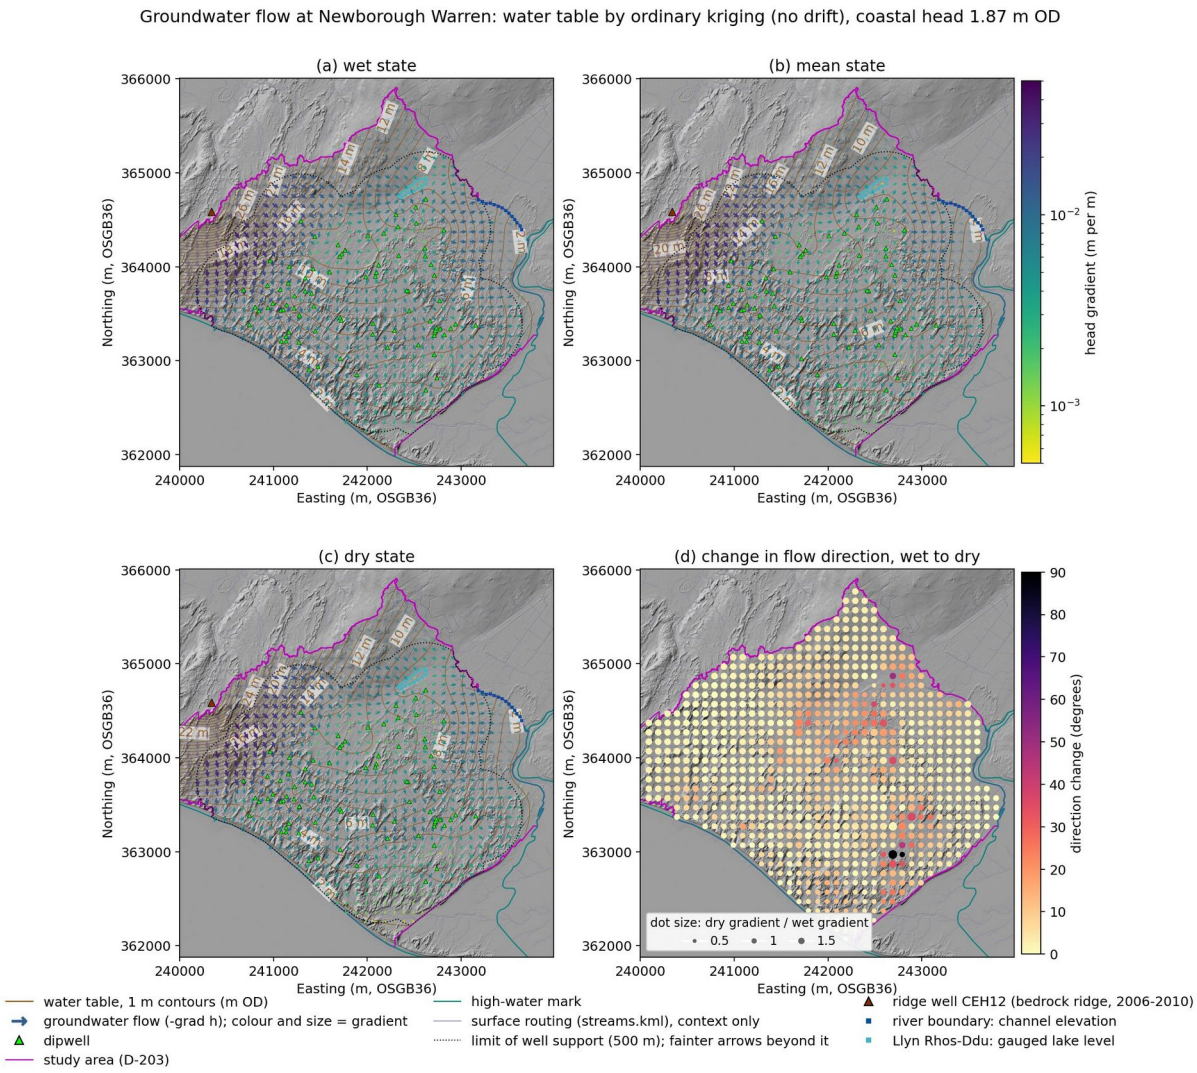


Figure 15. *The water table kriged from the dipwells in (a) the wet state (wettest tenth of months), (b) the mean state and (c) the dry state (driest tenth), with 1 m contours (m OD). Arrows show the negative head gradient, one per 100 m square, sized and coloured by the gradient; arrows more than 500 m from a dipwell are faint, because there the boundaries rather than the wells set them. (d) The change in flow direction from wet to dry. Background: LiDAR hillshade (© NRW & OS).*

4.10 Water-balance residual at the wells

The water-balance residual is modelled losses minus modelled recharge at each well, calculated from the fitted coefficients. It closes tightly, carries no coherent spatial structure and has a small systematic negative offset (Figure 16). Residuals span −0.0105 to +0.0080 m/month about a median of −0.0040 m/month. Of the 66 reference wells, 65 lie inside ±0.01 m/month and 51 are negative. Four of the five cluster medians are negative, from −0.0067 m/month at C4 Main Forest to −0.0010 m/month at C5 Coastal Forest. C1 Lake Edge is slightly positive (+0.0006 m/month). Residual magnitude is uncorrelated with position on either axis (Spearman ρ = −0.14 against Easting, p = 0.27; ρ = +0.16 against Northing, p = 0.20). The signed residual shows no gradient on either axis (ρ = +0.02 against Easting, p = 0.86; ρ = −0.10 against Northing, p = 0.40).

No well exceeds +0.02 m/month, and only fifteen are positive at all. The largest are NW3 (+0.0080) in C2 Dune, CEH23 (+0.0074) in C1 Lake Edge and NW6 (+0.0073) in C3 Western Residual. All three lie in the open dune or at the lake edge rather than at the ridge. No C4 Main Forest well is positive. The largest negative residual falls at the ridge-flank well CEH14 (−0.0105 m/month), with CEH32 (−0.0079 m/month) next. Two independent diagnostic tests on the monthly SSM residual both returned null results. One was a cross-correlation lag test for distance-dependent transport from the ridge; the other was a seasonal climatology test for unmodelled summer evaporative demand. No unmodelled flux is required to close the balance at any cluster.

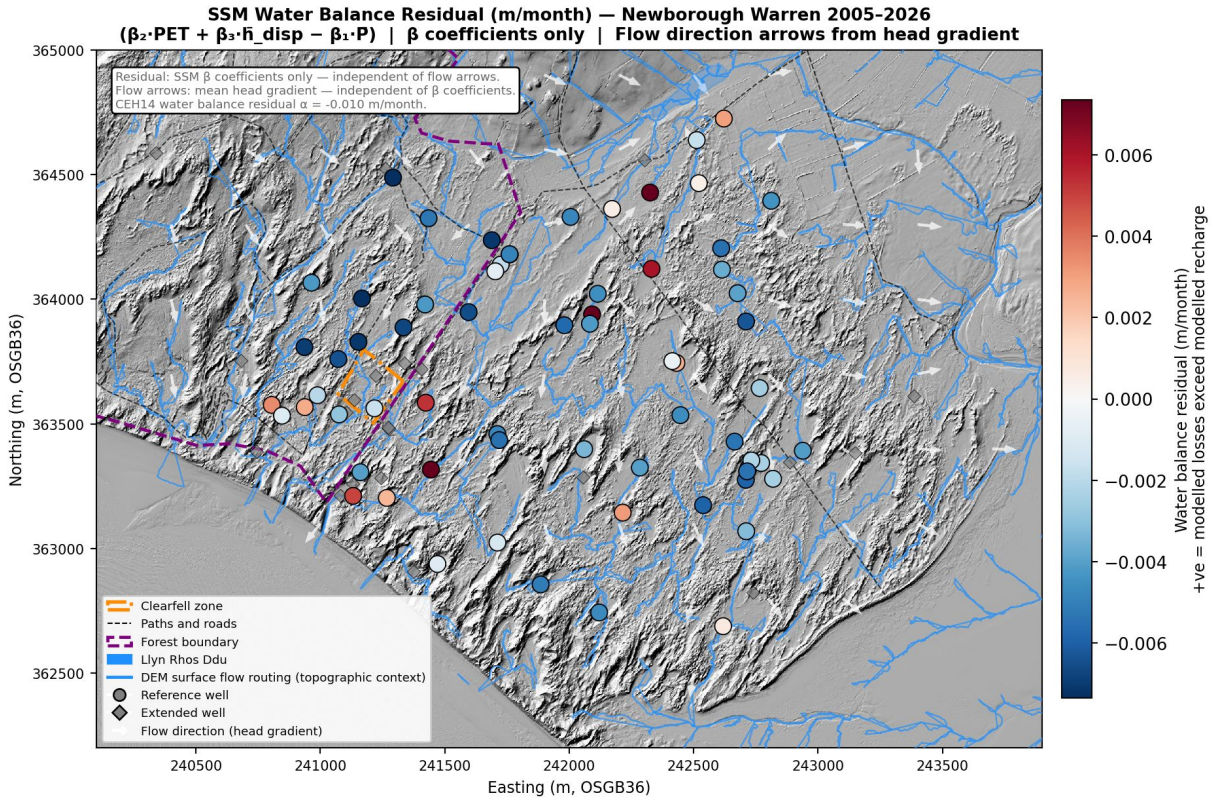


Figure 16. *SSM water-balance residual at the 66 reference wells. Positive (red) indicates modelled losses exceeding modelled recharge; negative (blue) indicates over-prediction. The residual closes tightly across the network, with a small systematic negative offset: 65 of 66 reference wells lie within ±0.01 m/month and 51 are negative. No well exceeds +0.02 m/month; the largest positive value is NW3 (+0.0080 m/month). Rainfall is gross, and climate means are taken over the same period as the mean heads. Residuals are shown at the wells alone: no interpolation between wells predicts a withheld well's residual appreciably better than the network mean, so no surface is drawn. The residual is a diagnostic of a lumped-parameter model, not the output of a calibrated continuous-flow simulation.*

4.11 Coastal-retreat gradient and the partition of long-term decline

4.11.1 The coastal gradient

The network-scale regression of water-table trend against distance from the eroding Caernarfon Bay shoreline selects a linear-with-cutoff (Dupuit–Forchheimer) form in the forest-free headline specification (Figure 17). Relative to the far field, the coast-edge deepening rate is −26.38 mm/yr, evaluated at 150 m from the shoreline (95% CI −33.23 to −19.53, well basis). The inland reach is about 900 m (620–1185), and the fitted far-field asymptote is −0.30 mm/yr. The rate is quoted at 150 m rather than at the shoreline because no well sits at the shoreline. The decay amplitude there, δ_0 = −31.28 ± 4.79 mm/yr, is an extrapolation beyond the network. The shoreline is also the one distance at which the two candidate decay forms disagree materially.

The forest-free and full-network fits agree closely, and the C3-only specification returns a consistent coast-edge rate. This indicates that the gradient is not an artefact of forest-cover inclusion. Model selection by AIC favours the exponential-decay form in the forest-free specification (ΔAIC = −6.0, exponential − linear-with-cutoff). The linear-with-cutoff form is adopted as the headline on physical grounds: it embodies the Dupuit–Forchheimer expectation of a finite inland reach, beyond which the migrating seaward boundary exerts no influence. It is also adopted because its far-field asymptote takes the same sign as the observed far-field trend, where the exponential's does not. Quoting the rate at 150 m makes that choice largely immaterial: the two forms agree to within 1.6 mm/yr there, against 9.3 mm/yr at the shoreline.

4.11.2 How robust the gradient is

The amplitude does not depend on how the forest is treated. The full network without a canopy control gives δ_0 = −31.58 mm/yr, the full network with a canopy covariate −32.26, and the Western Residual wells alone −28.28, against the forest-free headline of −31.28 (Figure 17a; a comparison of point estimates). The coast-proximal decline is therefore not an artefact of the plantation's own drawdown.

The amplitude does carry one well's leverage. Withholding CEH3, the second-nearest well to the shore at 176 m, moves δ_0 by 3.1 mm/yr and the 150 m rate by 2.4 mm/yr. That is 1.9 times the next-largest single-well shift (each well withheld in turn over the 61-well panel). Uncertainties on the coastal gradient are quoted on the well basis. The standard error is estimated from those refits, because consecutive monthly readings within a well are not independent observations. The fitted standard errors from the monthly rows are retained as the nominal precision of the model, not quoted as uncertainties. On that basis δ_0 is −31.28 ± 4.79 mm/yr, more than six standard errors from zero, and the leverage of CEH3 lies inside the error. The sensitivity is a property of a network with few wells inside 300 m rather than an instability of the fit. It is stated because a single near-shore well should not carry a headline unremarked.

4.11.3 What the far-field asymptote means

δ_0 and L_cg are identified by the distance dependence of the trend and are unaffected by the level. The asymptote is a different quantity. Two senses of "identified" bear on it, and both are worth stating. Statistically, it is separately identified from the cumulative-water-balance covariate that the panel uses to absorb spatially uniform climate forcing: the two columns are not collinear (variance inflation factor 1.01). Interpretively, it is not. A constant drift and the covariate's own trend contribution over the panel's span trade off against one another, so only their sum is recoverable. The asymptote is also unstable with respect to the fitting window. It runs from −0.30 to +23.96 mm/yr as the window start is moved with the well set held fixed. On either reading it is a window statistic rather than a rate, and no far-field climate rate is read from it here.

4.11.4 Seasonality

The gradient is itself seasonal, but the boundary effect it expresses is year-round. A season × distance interaction on the forest-free panel (10,909 observations) fits the coastal term as $\delta(d)\cdot t\cdot(1 + \gamma\cdot S)$, with S = 1 in Mar–May. It returns $\gamma = +0.309 \pm 0.060$ ($p < 0.001$) for the linear-with-cutoff form and $+0.304 \pm 0.059$ ($p < 0.001$) for the exponential. The coastal-retreat drift rate is thus about 31% steeper in Mar–May than over the rest of the year, at the same magnitude under both decay forms. γ scales the drift up in spring rather than creating it there: outside March–May the drift still runs at about three-quarters of its spring rate. The forested C5 wells, which are not in this panel, likewise decline in both the summer minimum and the spring mean (−32 and −24 mm/yr, balanced-basis slopes). The all-season fit reported here is therefore a season-averaged quantity, not a season-independent one.

4.11.5 How much of each cluster's decline is coastal

Applying the headline fit to each cluster's summer-minimum decline partitions the trend against a balanced observed basis (Table 8). That basis is the slope of the annual cross-well mean of the per-well summer minima. It is declared in the output itself, so that the basis travels with the table. The model has two parts: a coastal term, which depends on how close a cluster's wells are to the shore, and a non-coastal term that is the same at every cluster, +2.65 mm/yr (the trend the cumulative-water-balance covariate carries, +2.95, plus the far-field asymptote, −0.30). The unexplained remainder is the observed decline minus both. The robust result is the network curve itself, fitted across thousands of well-months of the open-dune network and insensitive to functional form, together with the C3 Western Residual decomposition.

Eleven of C3's twenty-two wells in the gradient subset lie within the fitted inland reach, including the four nearest the shore. Coastal retreat accounts for −7.0 mm/yr, most of C3's decline on that basis (93%; 87–102% across delete-one-well refits of the gradient, which move δ_0 and L_cg together). This holds although the cluster's mean distance (906 m) falls just beyond the reach. C2 has no well within the reach and carries no gradient component. C1 has one, CEH11 at 783 m, and carries −0.6 mm/yr. C4's mean distance falls marginally inside the reach, at 889 m against a reach of about 900 m, and three of its four wells lie within it. It carries a component of −4.8 mm/yr. That is 74% (58–91% across the same refits) of a balanced decline

(−6.5 mm/yr) that the record does not resolve; with the non-coastal +2.65 mm/yr the model gives −2.2 mm/yr in total. The reach's confidence interval runs from 620 to 1185 m, and the reach lengthens further under alternative climate covariates.

Against balanced declines of −11.0 and −14.6 mm/yr at C1 and C2, the modelled totals are +2.1 and +2.6 mm/yr. These totals are the trend contribution the cumulative-water-balance covariate carries plus the fitted far-field asymptote, with C1's small coastal term. The whole of their deepening, and a little more, is therefore left unexplained rather than attributed to climate. At C3 the unexplained remainder is −3.2 mm/yr, and at C4 −4.3 mm/yr.

This decomposition identifies coastal retreat as a distinct, spatially structured forcing on long-term summer minima. A climate-only attribution would fold it into an apparent network-wide decline. The coastal contribution declines linearly across the reach and is zero beyond it, so it is a spatially structured forcing rather than a uniform one. The spatially uniform component is a window statistic rather than a rate. No distance at which the two are equal follows from this fit.

4.11.6 An independent check along a transect

An orthogonal, longitudinal test corroborates the cross-sectional gradient. It follows a shore-normal transect from a coastal anchor (CEH22) to an inland anchor beyond the fitted reach (NW4). Along it, the spring (MAM) coast-to-inland head difference deepens at −27.8 mm/yr (bootstrap 95% CI −36.1 to −19.3; AR(1)-corrected; 2010–2023). This is consistent with the modelled coast-edge rate of −26.38 mm/yr. The coastal anchor sits at 147 m, so the comparison is made at essentially the distance at which the modelled rate is quoted. Referencing the coastal record to the erosion-free inland anchor removes common-mode climate without a model. The coastal-end drawdown then orders monotonically through time (Spearman ρ improving from −0.21 for the raw record to −0.79 after referencing; Figure 17c). A coast-to-inland difference that grows through time is the signature of a migrating seaward boundary, not of a static geometric offset or a climate wobble. It provides an observational counterpart to the cross-sectional fit on the same record.

4.11.7 C5 Coastal Forest

The C5 Coastal Forest row is reported with a caveat rather than as a firmly quantified cluster attribution. Four C5 wells enter the gradient subset after the clearfell-zone exclusion and the requirement for a valid summer-minimum slope. They are the reference wells CEH3 at 176 m from the coast, NW9 at 419 m and NW8 at 521 m, and the extended well NW8b at 524 m (mean 410 m). The three reference wells decline at −32.0, −32.7 and −59.2 mm/yr. Each is statistically significant ($p = 0.0002$, 0.002 and 0.005), and together they are much the steepest group in the network. NW8b, on an eleven-year record, shows no significant trend (+8.5 mm/yr, $p = 0.65$).

The gradient attributes −17.1 mm/yr across the four. That is 53% of the cluster's balanced basis of −32.0 mm/yr; with the non-coastal +2.65 mm/yr, −17.6 mm/yr is left unexplained. The per-well mean (−28.9 mm/yr) and the Script-14 centroid (−29.1 mm/yr) lie close to that basis. The share is sensitive to which wells the cluster contains. Before NW8b was assigned to C5, the

same decomposition left −0.7 mm/yr unexplained (Hollingham, 2026d). The percentage should therefore be read as indicating that coastal retreat accounts for about half of C5's decline, not as a calibrated share.

The forest-free curve is fitted on the open-dune wells and excludes the forested clusters entirely, so none of the four can bias the fit. Nor does the choice of functional form bear on the reading. Across their distance range the linear-with-cutoff form predicts deepening from −25 mm/yr at CEH3 to −13 mm/yr at NW8, and the exponential is of the same order throughout. The cluster is therefore treated as corroborating the gradient's direction and magnitude rather than as an independent cluster-level estimate. It is cross-referenced to the substrate and coastal-position interpretation of C5 in Section 4.8.

Table 8. *Decomposition of cluster summer-minimum decline (mm/yr) against a balanced observed basis. The observed column is the slope of the annual cross-well mean of the per-well summer minima over the gradient-analysis subset — one regression on one annual series, rather than an average of per-well fits taken over unequal record windows, which reports a composition of record lengths in place of a rate. The climate + offset column is the trend contribution the cumulative-water-balance covariate carries over the fitted span plus the fitted far-field offset c_far; the asymptote is a window statistic rather than a rate, so this column is not a measured climate background. Well counts here are the gradient-analysis subset (clearfell-zone wells dropped; wells require a valid summer-minimum slope), not the cluster sizes of Table 1. Components may not sum exactly to the observed basis because of rounding.*

Cluster	mean dist. (m)	observed	coastal-retreat	climate + offset	unexplained	coastal %
C1 Lake Edge	1819	−11.0	−0.6	+2.6	−13.0	5
C2 Dune	1565	−14.6	0	+2.6	−17.2	0
C3 Western Residual	906	−7.5	−7.0	+2.6	−3.2	93
C4 Main Forest (n = 4)	889	−6.5	−4.8	+2.6	−4.3	74
C5 Coastal Forest (n = 4) †	410	−32.0	−17.1	+2.6	−17.6	53

† C5 rests on four wells in the gradient subset, three reference wells and the extended well NW8b, and lies closest to the shore; the share is sensitive to which wells the cluster contains. Reported for completeness, not as a calibrated attribution (see text). The Script-14 cluster-centroid trend, retained in the output as a context column, is −29.1 mm/yr against a balanced basis of −32.0 mm/yr. The coastal term of every row is the mean over the cluster's wells of the capped gradient at each well's own distance; the distance column is the cluster mean, given for orientation.

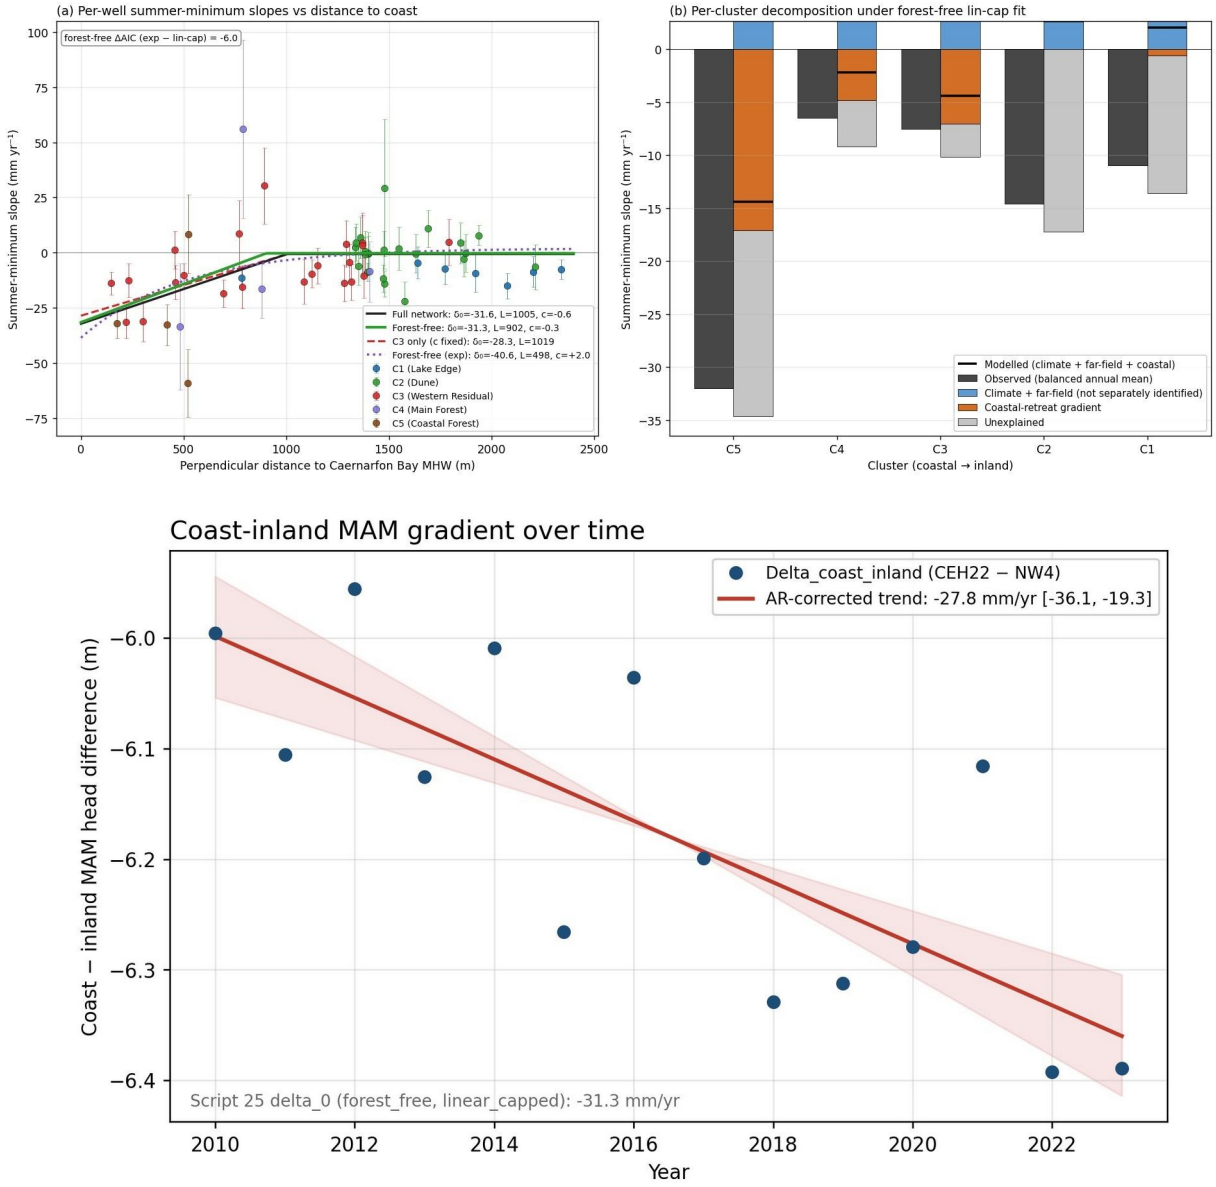


Figure 17. *Coastal-retreat gradient. (a) Per-well summer-minimum slope versus distance to the eroding Caernarfon Bay shoreline, with the full-network, forest-free and C3-only linear-with-cutoff fits and the forest-free exponential fit overlaid; b) per-cluster decomposition of the observed decline into climate, coastal-retreat gradient and residual; (c) observational corroboration — the spring (MAM) coast-to-inland head difference (coastal anchor CEH22 minus inland anchor NW4) against year, with the AR(1)-corrected trend of −27.8 mm/yr; this transect subset is distinct from the whole-network fit in (a) and (b).*

5. Discussion

5.1 Geology, land cover and the cluster structure

The five hydrogeological clusters, resolved by hierarchical clustering of hydrograph dynamics, map onto the documented geological structure of the site. In the dendrogram, the top Ward merge (≈0.99) separates C4 Main Forest from the other 57 wells. The east–west division follows at ≈0.61. It separates the responsive eastern open dune (C1 Lake Edge, C2 Dune) from the buffered west (C3 Western Residual with C5 Coastal Forest). That contrast reflects the difference in subsurface stratigraphy described by Stratford et al. (2007). In the east, shallow sand over till and estuarine deposits constrains storage and drives flashy water-table responses. In the south-west, a deeper sand column supports a more buffered aquifer.

The mechanistic coefficients recover this divide independently of the clustering that defined it. They also resolve a gradient running through the open dune, rather than a simple binary. C1 Lake Edge is one extreme. It has the highest recharge sensitivity in the network ($\beta_1 = 4.58$), the fastest head-dependent drainage ($-\beta_3 = 0.089$) and the shortest drainage half-life in the network ($t_{½} \approx 8.0$ months). This is a fill-and-drain regime, in which a shallow base promotes rapid lateral exchange directly with Llyn Rhos-Ddu. C2 Dune is larger than C1 and is the most representative open-dune unit (19 wells). It is the open-dune archetype between the end-members. It has a still-high recharge sensitivity ($\beta_1 = 3.90$) and a drainage-dominated loss partition (68% of losses), but a slower drainage decay ($-\beta_3 = 0.064$) and a longer half-life ($t_{½} \approx 10.7$ months) than the lake-coupled C1. This reflects open dune without a direct lake connection. C3 Western Residual is the buffered opposite. It has the lowest recharge sensitivity of the open clusters ($\beta_1 = 3.76$), the slowest open-system drainage ($-\beta_3 = 0.062$) and the slowest open-dune drainage ($t_{½} \approx 11.6$ months at the project datum). Here the deep aeolian sand attenuates both the rainfall pulse and its release. The forested interior fills slowly and drains most sluggishly of all. C4 Main Forest is the cluster geographically nearest the rock ridge, and the metamorphic bedrock outcrop forms its northern margin. That setting is reflected in the impeded drainage and the elevated atmospheric-draw coefficient (Section 5.2). It is revisited in the residual-field discussion (Section 5.4).

This gradient also bears on how C3 is best understood. Its membership is the most diffuse in the network. The great majority of its wells are gradationally assigned: they correlate almost as strongly with the C2 Dune centroid and with the coastal forest as with their own. It is also the open-dune cluster closest to both the plantation margin and the eroding Caernarfon Bay shoreline. C3 is therefore better read as the transitional, more strongly buffered south-western continuation of the open dune than as a discrete western block. It has a coherent central signature (the longest open-dune half-life and the lowest open recharge sensitivity), but its boundaries grade into both the dune and the forest.

One parsimonious mechanism is consistent with the surface geology, though not directly observed here. A single estuarine-clay base crops out on the higher eastern and southern ground, where a thin overlying sand column produces the flashy C2 response. It dips beneath

the surface toward Caernarfon Bay, where a thickening sand column buffers the C3 response. This would make C2 and C3 one dipping system rather than two substrate blocks.

The gradient within C3 is consistent with this. Across the 21 C3 wells that carry a coastal distance, recharge sensitivity β_1 and the drainage coefficient β_3 both rise inland ($r = +0.77$, $p < 0.001$; $r = +0.58$, $p = 0.006$). The water-table-fluctuation specific yield falls inland, from 0.40 at the south-western margin (CEH21) to 0.22 at the north-east (D38; $r = -0.80$, $p < 0.001$). The response thus becomes more buffered toward the coast, which is the opposite of the flashier response a discharge boundary alone would produce. It grades toward the open dune at the cluster's inland edge. Because β_1 is recharge divided by specific yield, a single control accounts for both coefficient gradients: a south-westward rise in near-surface specific yield.

The storage evidence points the same way and is independent of the coefficients. Within the open dune, the water-table-fluctuation specific yields rise on a coherent planar trend toward the south-west (Section 4.5). That is the direction in which a dipping base would thicken the sand column. Coastal proximity, the plantation margin and the inferred clay dip all covary along the same south-western axis. This evidence therefore supports the transitional reading without isolating the clay dip as its sole cause. A cored or geophysical transect would be needed to test it directly (Section 5.7). Figures 18 and 19 map the two reaches responsible: the forest-interception drawdown propagating inland from the plantation, and the coastal-retreat drawdown extending inland from the eroding front. They place the C3 wells within both reaches, while the C2 wells lie largely beyond them.

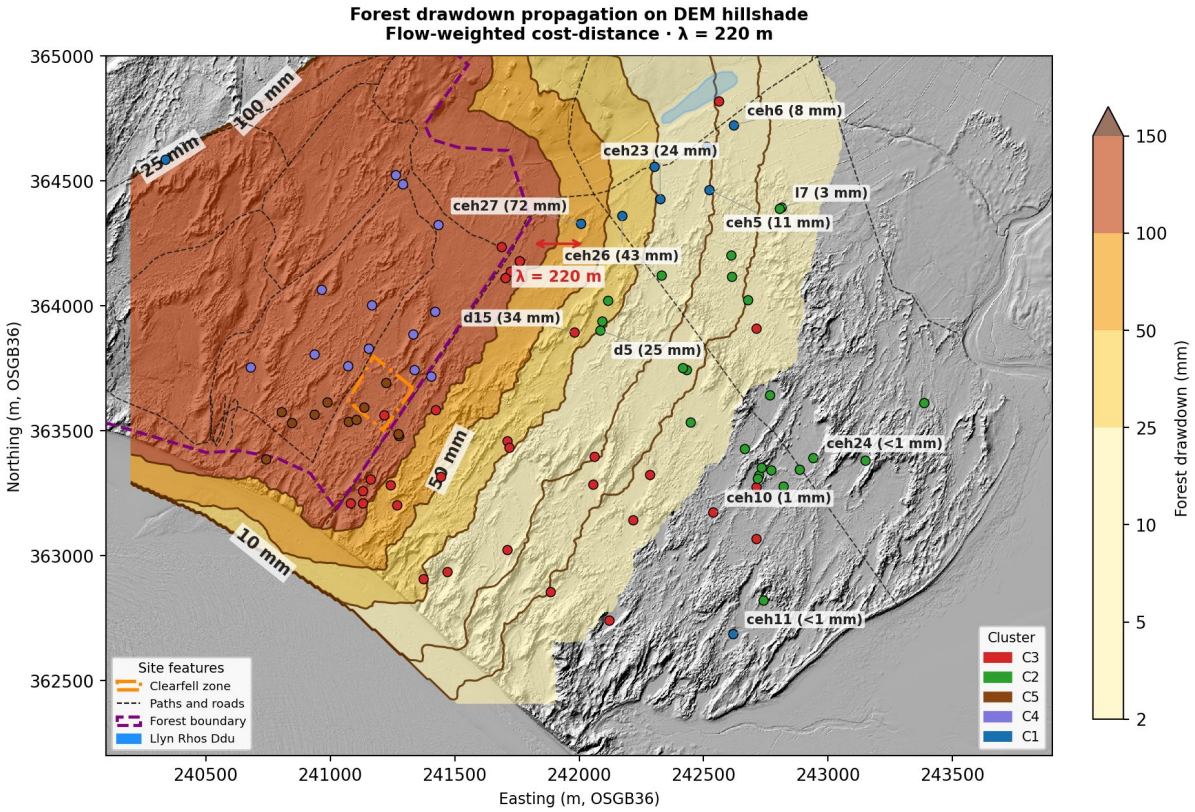


Figure 18. *Illustrative forest reach. Forest-interception drawdown propagated inland from the plantation boundary by flow-weighted cost-distance (decay length λ ≈ 220 m, from C3's specific yield and drainage coefficient combined with an assumed hydraulic diffusivity — a literature conductivity, K = 6 m/day (Betson et al., 2002), and a nominal 5 m saturated thickness, the latter poorly constrained and plausibly larger; assumed edge deficit). The reach is a scaling estimate rather than a calibrated length: across the range reported for dune sand (K = 2–20 m/day) λ spans roughly 130–410 m, and across a 3–8 m saturated thickness 170–280 m, so it is stated here to the nearest tens of metres. The mapped pattern depends only on the reach being of order 200 m rather than of order 20 m or 2 km, which holds across that whole range, over the DEM hillshade with per-well drawdown values labelled. The C3 wells fall within the forest reach while the C2 wells lie largely beyond it — drawdown falls from tens of millimetres near the plantation to under 1 mm at the easternmost C2 wells; the forest boundary, clearfell zone and Llyn Rhos-Ddu are marked. Drawdown magnitudes are illustrative of the reach rather than measured values; because the model represents canopy interception alone, its heads are a lower bound on the magnitude of the forest effect, and λ an upper bound on how far beyond the canopy that effect reaches.*

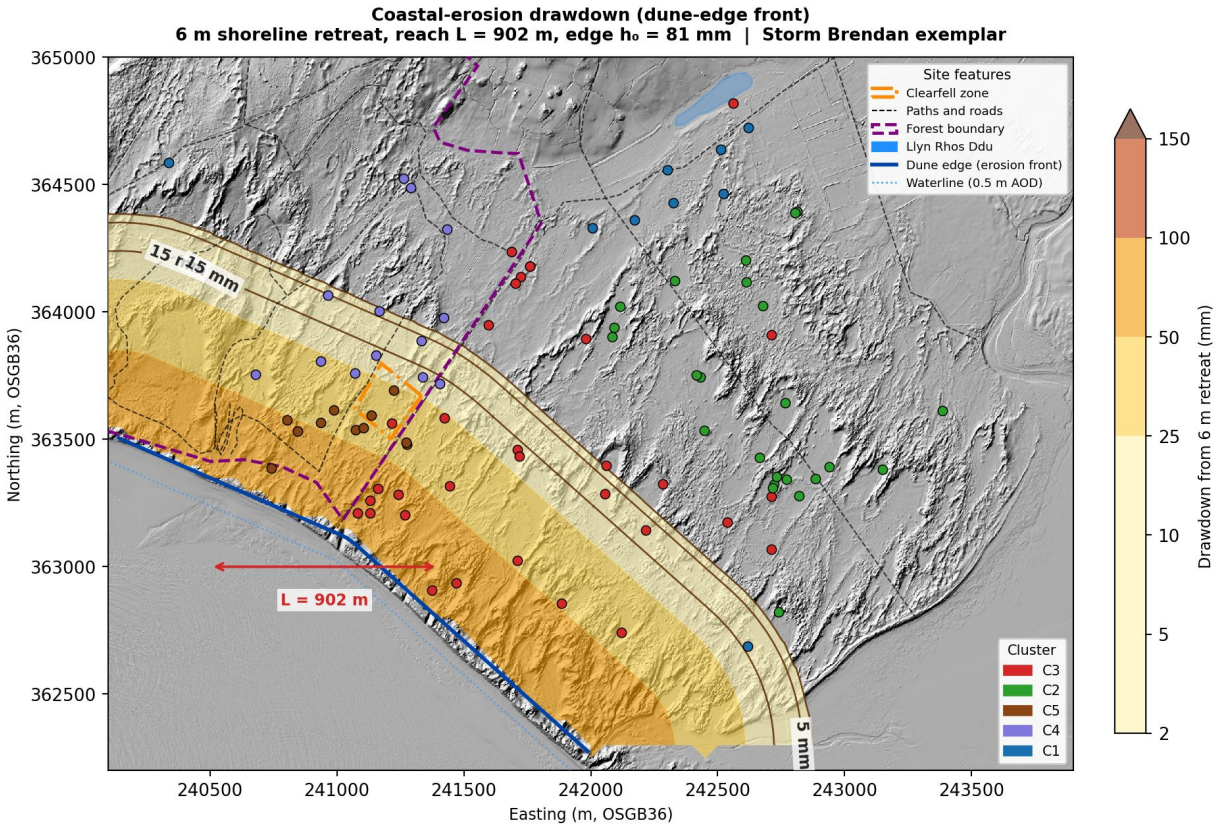


Figure 19. *Episodic coastal-retreat reach. Water-table drawdown contoured inland from the eroding dune-edge front for a single Storm Brendan-scale retreat event (≈6 m at the front, early 2020), over the fitted inland reach (L_cg ≈ 900 m; §4.11). Coastal retreat at this shore is episodic, concentrated in storm events; the per-event edge drawdown is scaled from the chronic gradient through the retreat rate measured over the same period as that gradient (E_c = 2.32 m/yr, 2006–2026, from shoreline positions digitized in this study). The C3 wells fall within the reach while the C2 wells lie largely east of it; the forest boundary and Llyn Rhos-Ddu are marked.*

The clustering (Figures 4 and 5) and the mechanistic coefficients (Figures 11–13 and Table 1) show the same structure. This is consistent with a partition that reflects the aquifer rather than the linkage method. The drainage decay half-life (Table 6; Figure 14) orders the clusters the same way. C1 Lake Edge drains fastest and C4 Main Forest most slowly. C5 is distinct from C4 despite sharing the same canopy.

5.2 Substrate versus surface control: interpreting the coefficients

The two forested clusters are inferred to share the deep sandy substrate of the western block, and they carry the same Corsican pine canopy. Yet they segregate into distinct units. Resolving why is central to interpreting what the SSM coefficients measure.

Within the forest zone, elevation alone explains 98% of the variance in atmospheric draw (β_2 versus elevation, r = 0.989). C4 and C5 separate completely on this axis and on mean elevation (10.6 m versus 5.6 m). C4, on the elevated ridge flank, has a deeper water table and unsaturated zone, together with weaker drainage. We interpret this deeper profile as permitting stronger atmospheric draw. C5, on the lower coastal fringe, has a shallower water

table and faster drainage. The measured contrast in water-table depth and coefficient behaviour is clear.

That interpretation has independent measured support at this site. Freeman (2008) recorded the water table beneath the Newborough plantation approximately 1.07 m below the open Warren (ANOVA $p < 0.001$). She also recorded mean volumetric soil moisture there at roughly half the open-dune value. She attributed the moisture contrast to the loss of capillary contact between the upper soil layers and the water table where the water table is deeper. The forest zone therefore carries an unsaturated zone that is both thicker and drier, consistent with the coefficient contrast we recover.

The substrate explanation differs in confidence between the two clusters. For C4, the bedrock of the ridge outcrops within the forest, so the nature of the rock underlying the cluster is established. Its depth and extent beneath C4 are not known without coring. But the impeded drainage there is consistent with a shallow, irregular bedrock base close beneath the wells, an inference supported by the proximity of the outcropping ridge. The storage evidence within the plantation supports it directly. Across the fourteen forested wells, the specific yield falls with elevation ($r = -0.93$, $p < 0.001$) and rises with distance from the ridge ($r = +0.81$, $p < 0.001$). This is the pattern expected where a shallowing bedrock base thins the drainable sand column toward the outcrop. The relationship is robust to excluding the two anomalous ridge-flank wells ($r = +0.79$ against ridge distance with both removed). It is present in the uncorrected estimates as well as the corrected ones ($r = +0.70$, $p = 0.005$), so it is not an artefact of the canopy correction. For C5, the proposed more-uniform, deeper coastal sand is the weaker inference, resting on the coefficients and topography alone.

On that basis, the split is best read as a substrate-and-topographic-position discontinuity beneath a constant canopy, rather than an arbitrary division of a continuous population. The forest-interception drawdown footprint (Figure 18) occupies the same corner. The high atmospheric draw, the positive water-balance residual and the modelled drawdown reach therefore together describe one coherent forest signature. All three, however, are anchored to the same plantation-and-ridge zone, so their spatial agreement is mutual consistency rather than independent corroboration.

This finding clarifies a potential misreading of the cluster contrasts. The forested clusters carry the highest LCSC values in the network, which might be read as low aquifer drainability. But LCSC is a recharge-sensitivity descriptor, not a storage parameter. The fitted β_1 absorbs canopy interception alongside the true storage response. Part of the forest clusters' elevated LCSC therefore reflects the suppression of throughfall by the canopy, rather than a difference in the aquifer itself.

The water-table-fluctuation specific yields bear on this directly. The uncorrected forest medians exceed the open-dune range. This is because the method attributes the full rainfall total to recharge, when in reality the canopy intercepts roughly a quarter of it. Once corrected for interception, C4's forest median (0.261) falls within the open-dune range (0.210–0.325). C5's (0.330) sits just above its upper edge, though C5's estimate is constrained by the physical-

plausibility limit and is only weakly corroborative. The interpolated specific-yield surface shows no discontinuity at the cluster margins.

The effective storage inferred from the water-table response therefore varies smoothly and continuously across the site, stepping nowhere at the cluster margins. This is consistent with, though on these data not proof of, a well-mixed aeolian depositional history. The C4 specific yield and the continuous Sy surface, rather than the ceiling-bound C5 estimate, carry the argument that the forest clusters' distinctive coefficient signature derives largely from a surface boundary condition (canopy interception) rather than from a different subsurface architecture.

A depth-dependent extension of the model lets atmospheric draw decline as the capillary fringe retreats from the root zone with increasing water-table depth. It indicates that a real depth-coupling effect is present at every cluster. The effect is strongest where the water table is shallowest (C1) and weakest where it is deepest (C4). Notably, C4 shows the weakest depth coupling despite carrying the highest atmospheric-draw coefficient. This is consistent with its high β_2 reflecting low specific yield rather than a depth-dependent evapotranspiration pathway. The standard fixed-β_2 model absorbs the coupling as an effective depth-averaged parameter. The depth coupling and the lateral-drainage geometry together contribute to the within-forest β_2 variation, alongside substrate. But the dominant control on the cluster-level contrast remains topographic position and substrate rather than canopy. These within- and cross-cluster comparisons rest on the ordering of β_2 rather than on its absolute magnitude. The fitted β_2 values are OLS estimates biased somewhat high in absolute terms. They are reliable as rankings rather than as quantitative atmospheric-draw rates.

5.3 Recharge, loss and the water balance

Decomposed under identical climate forcing, the clusters reveal markedly different recharge sensitivities and loss partitions. At the same time, the decomposition closes the monthly water balance to within 1.8% of total losses at every cluster. Recharge declines systematically from the Lake Edge through the open dune to the forested clusters, and total losses track recharge closely. The partition between drainage and atmospheric draw, however, varies sharply. Drainage dominates at C1 (85% of losses at the project datum), reflecting rapid lateral exchange with the lake. Atmospheric draw dominates at C4 (76%). There, the suppressed recharge pulse never lifts the water table far enough above equilibrium to generate a strong drainage response, and the irregular substrate impedes what drainage there is.

The suppression has two components. Interception removes a fraction of rainfall at the crown. In addition, the profile beneath the canopy is measurably drier, so a larger storage deficit must be satisfied before the remainder reaches the water table; Freeman (2008) recorded roughly half the open-dune volumetric soil moisture beneath the plantation. Her felled clearing plot carried a water table statistically indistinguishable from the open Warren. This indicates that the second component is a consequence of the canopy rather than an independent property of the forest zone.

The displacement formulation is what makes this contrast resolvable. Fitted on depth below ground alone, the drainage coefficient is negative at two of the five clusters. From a datum of 1.6 m downward, it is positive and significant at all five (Hollingham, 2026c).

Two cautions attach to the recharge term. First, the model is fitted without an intercept. In an open-boundary system the recharge coefficient can therefore absorb unmeasured boundary fluxes (Young, 2011; Knotters and van Walsum, 1997). $\beta_1 \cdot P$ should not be read as a direct measure of local precipitation reaching the water table. Second, the atmospheric-draw coefficient is an effective parameter. Thornthwaite PET is an energy-based demand index, and it does not define a reference surface in the manner of Penman–Monteith. β_2 therefore folds transpiration, interception, specific yield and depth coupling into a single scaling between atmospheric demand and water-table response, and a formal crop-coefficient decomposition does not apply. Both cautions are consistent with the substrate interpretation above. The coefficients are diagnostic of behaviour. Their physical decomposition requires the independent specific-yield and depth-coupling evidence assembled here.

5.4 Boundary conditions and the water-balance residual field

The near-closure of the water balance at the cluster mean also holds in space. The residual at the wells is modelled losses minus modelled recharge, from the fitted coefficients. It lies within ±0.01 m/month at 65 of the 66 reference wells and carries no gradient across the site. Its cluster medians span −0.007 to +0.001 m/month. A positive residual would mark a location where measured rainfall-driven recharge does not fully account for the observed water level, implying an unmodelled input. No well carries such a residual above +0.02 m/month. The fifteen positive residuals are all below +0.01 m/month, and they are largest in the open dune and at the lake edge (NW3, CEH23, NW6) rather than at the ridge. The largest negative residual falls at the ridge-flank well CEH14. On this evidence, no ridge-derived lateral input is required to close the balance, and the residual field provides no positive support for one.

Two further diagnostics constrain what the record could have detected had such an input been present. Two independent diagnostic tests both returned null results: a cross-correlation lag analysis for distance-dependent transport from the ridge, and a seasonal climatology analysis for unmodelled summer evaporative demand. The cross-correlation null is itself qualified by sampling frequency. Monthly observations are an order of magnitude coarser than the days-to-weeks transit times expected for fracture-flow input from a metamorphic bedrock ridge. A real ridge-derived lag could therefore be invisible at this resolution. The null result bears on what the present record can resolve, not on whether the mechanism is operative.

The seasonal phase of the residual narrows the field further. The residual peaks in winter and early spring at the great majority of wells, and at none in summer. That is inconsistent with any vertical-flux parameterization error: systematically underestimated summer PET, an incomplete canopy-interception correction, or any under-resolved summer evaporative demand. All of these would peak in summer, when those fluxes are active. The phase is consistent instead with either a winter-phased nonlinear recharge mechanism or ridge-derived lateral input arriving during the months of greatest rainfall.

A cluster-stratified version of the same climatology shows mean winter-minus-summer residual contrasts of +8.6 mm in C3 ($p<0.001$), +5.4 mm in C4 ($p=0.068$, marginal) and −0.3 mm in C5 (not significant). The C3 > C4 > C5 ordering and the magnitude of the C3 signal together rule out canopy-interception over-estimation as the sole driver of the winter–summer contrast, since the largest signal sits in a cluster with no canopy.

The residual pattern is real and reproducible, but its mechanistic attribution is unresolved. This is the appropriate stance for a lumped-parameter framework. The residual field is a structural diagnostic of where the three-term balance is insufficient. It is not a quantified map of lateral flux, because it also absorbs nonlinear soil-moisture dynamics, coefficient sensitivity and any systematic PET bias.

5.5 Long-term forcing: coastal retreat and the partition of decline

The network-scale gradient regression isolates a forcing that a climate-only analysis would misattribute. Fitted across the open-dune network, it resolves a coast-edge water-table deepening of roughly 26 mm/yr at 150 m from the shoreline. The deepening declines to zero over an inland reach of about 900 m, and of at most about 1 km under alternative climate covariates.

What that decomposition licenses is the removal of the coastal term, not the recovery of a climate rate. The fitted far-field asymptote is a window statistic rather than a rate: it moves across a range of some 24 mm/yr as the fitting window is varied with the well set held fixed. No clean climate baseline therefore follows from the fit, and none is offered here. An attribution analysis at this site must estimate its own baseline. What it should take from the decomposition is that a spatially structured coastal forcing exists at the western shore, and that this forcing has to be separated first rather than folded into an apparent network-wide decline.

The result is robust to functional form at the network scale and to forest-cover inclusion. It identifies coastal retreat as a spatially structured forcing distinct from climate. The gradient component is carried by the wells within the reach, so it is largest at C5, where coastal retreat accounts for about half of the observed summer-minimum decline. It accounts for most of the smaller decline at C3 (87–102% across delete-one-well refits), half of whose wells lie within the reach. C4 carries 74% of a decline the record does not resolve, C1 a negligible component and C2 none.

The water table takes a median of 5.1 months to return to its mean after a disturbance, measured without a model from the autocorrelation of the deseasonalized level (Hollingham, 2026c). This far exceeds the duration of individual storm-driven retreat events. The aquifer therefore functions as a low-pass filter on the shoreline forcing: it smooths discrete event-scale shocks into a chronic gradient. This supports the steady-state Dupuit–Forchheimer regression as the appropriate kinematic form.

The chronic gradient resolved here is the time-averaged expression of a retreat process that, at this shore, is episodic rather than steady (Pye & Blott, 2024). Shoreline loss is concentrated in storm events. The early-2020 Storm Brendan, a discrete retreat of order several metres at the

dune front, is a recent example. It is set against a mean of E_c = 2.32 m/yr, measured here from shoreline positions digitized across 2006–2026. That mean is the rate used to scale the gradient, because it spans the period over which the gradient is fitted. Reported rates for this frontage vary with window, extent and method. Pye & Blott (2024) record up to 16.3 m at the southern Twyni Penrhos profile over 2013–2022, about 1.8 m/yr at the most active point. A contemporary account of the 2014–2020 window reports roughly 50 m of loss concentrated in storm events (Forgrave, 2020).

Figure 19 maps the inland drawdown reach implied by a single Brendan-scale event. The steady-state regression averages many such events into the per-year deepening rate. The gradient and the episodic map are therefore two views of the same forcing, not competing estimates. Consistent with this, the coast-to-inland head difference is observed to deepen through the record at −27.8 mm/yr once common-mode climate is removed model-free (Section 4.11). A growing gradient is the temporal signature the steady-state regression predicts, and it distinguishes a migrating boundary from a static offset. On its own, though, it cannot separate erosion from a slowly varying substrate. The cored shore-normal transect (Section 5.7) remains the mechanism-resolving test.

Figure 20 draws the coastal forcing along that reach, largest at the shore and declining to nothing across the ~900 m span. The fitted far-field constant is a window statistic, not a site-wide climate background, and is not drawn as a curve. The same hydraulic diffusivity sets the equilibration time for a boundary perturbation to propagate across this reach, and that time is of order years to decades (Gerla, 2019). The coastal signal therefore registers as a chronic, quasi-steady gradient rather than a response to any single retreat event. A rising sea level, moreover, raises the water table at the coast, and the rise decays inland over a comparable reach. The fitted coast-edge gradient is thus the net of two opposing coastal processes. Sea-level rise of approximately 4 mm/yr (UKCP18 north-Wales near-term central estimate) works against the erosion-driven drawdown. This makes the fitted δ_0 a conservative estimate of the erosion component.

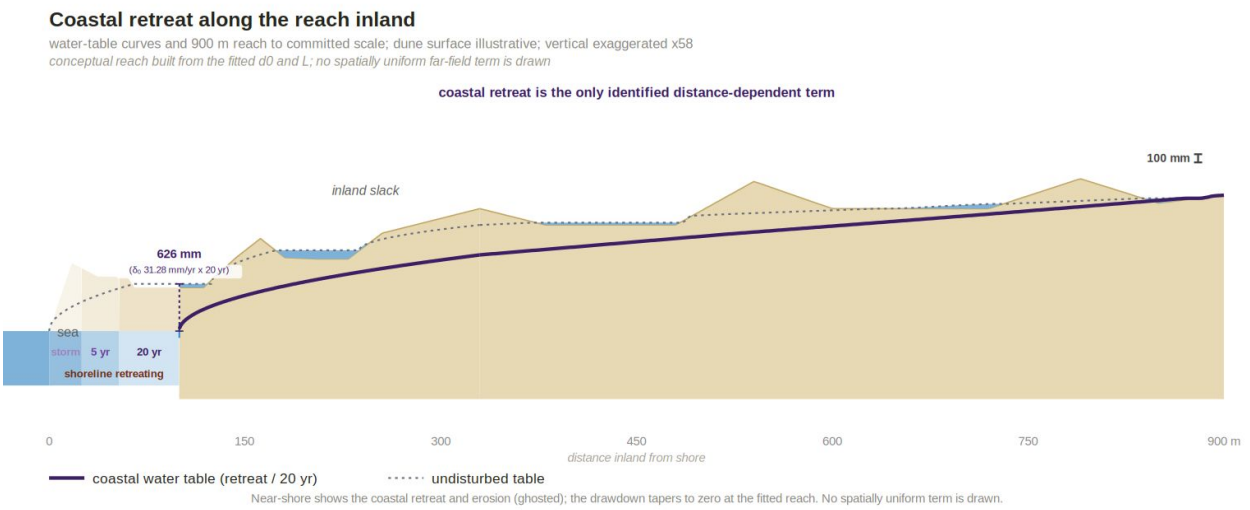


Figure 20. *Conceptual reach of coastal retreat and the far-field term. Long-section from the eroding Caernarfon Bay shoreline inland to 900 m, drawn so that the fitted inland reach L_cg ≈ 900 m (§4.11) falls at the right-hand edge; the dune surface and its slack are illustrative, the vertical scale is exaggerated ×58, and a 100 mm bar is given for reference. The coastal water table (solid) is the fitted coast-edge trend accumulated over a twenty-year horizon, declining linearly to zero at the fitted reach; the undisturbed table (dotted) is the same surface without it. At 150 m, where the headline rate is quoted, the twenty-year accumulation is 528 mm. The curve is drawn in to the shoreline, where the fitted amplitude δ_0 = 31.28 mm/yr gives 626 mm; that value is an extrapolation beyond the nearest well, which sits at 147 m. The fitted far-field constant c_far (−0.3 mm/yr, ≈ −6 mm accumulated over the same horizon) is a window statistic rather than a rate (§4.11) and is not drawn: over this reach coastal retreat is the only identified distance-dependent term. The ghosted wedges at the shore mark the shoreline positions implied by a single 6 m storm event and by five- and twenty-year accumulations of the mean retreat rate. The section is a schematic built from the fitted parameters rather than a modelled water-table surface. Adapted from the technical report (Hollingham, 2026d).*

This forcing also reframes the C5 Coastal Forest anomaly. C5 shows the steepest summer-minimum decline in the network, more than twice that of any other cluster. All clusters share the same climate forcing, so this argues against a purely climatic explanation. The position analysis (Section 5.2) and the coastal-gradient analysis converge on the same explanation. C5 occupies the lowest, most coast-proximal position in the network. There, an inferred thin saturated thickness would amplify the head response to any forcing, and the eroding shoreline lies closest.

Four C5 wells enter the gradient subset. One of them, NW9 (419 m from the coast), shows a decline of −32.7 mm/yr, among the steepest and most significant in the network. At that position the fitted gradient predicts a coastal-retreat deepening of roughly 15 mm/yr under either functional form. Because NW9 is a forested well, it is excluded from the open-dune (forest-free) regression. It therefore functions as an out-of-sample sentinel: its observed decline tests the fitted gradient at a near-coast position without contributing to it. The observed −32.7 mm/yr stands against a fitted coastal contribution of about −15 mm/yr. The residual is consistent with the site-wide decline and with the substrate amplification developed in Section 5.2. This agreement further corroborates the gradient at a near-coast position, complementing the shore-normal transect (Section 4.11).

A single well cannot, however, carry a quantified cluster attribution. The C5 figure should be treated as corroborating the gradient's direction and magnitude, not as an independent estimate. Even so, the substrate, position and gradient evidence converge, and this makes coastal proximity the leading explanation for C5's exceptional trajectory. It does not exclude a contribution from the maturing canopy, which this record cannot separate from the other two. The companion paper develops the management implication. Here the point is characterization: the framework distinguishes a coast-proximal, structurally amplified decline from the climate background. That distinction bears directly on how network-wide summer-minimum trends should be attributed.

One boundary mechanism operates alongside the retreat gradient at the most coast-proximal wells. The discharging foreshore is a tidal boundary, so near-coast water tables couple to a constant-head sea boundary. This buffers their short-term variability and holds their mean toward the shoreline. The tidal coupling and the multi-year shoreline retreat are complementary rather than competing. The former smooths event-scale fluctuation at a fixed boundary position. The latter migrates that boundary landward over years. The C5 decline is read against the slower retreat process. The tidal coupling is noted here as the boundary condition that makes near-coast records intrinsically less climate-legible, not as an alternative explanation for the trend.

The decline observed away from the coast, in the clusters that carry no gradient component or a negligible one (C2 and C1), is a candidate for further mechanistic work. At least two non-exclusive climate drivers are candidates. What the gradient regression cannot supply is the quantity those drivers would be fitted to. The panel absorbs the spatially uniform forcing through the cumulative-water-balance covariate. A constant offset trades off exactly against the trend contribution that covariate carries, so no residual far-field rate is isolated here (Section 4.11). The mechanisms below are accordingly motivated by the decline the network records, not by a fitted background. Climate forcing, separately measured over the station record, is the leading candidate for that decline.

The horizon of the decline is also bounded. Ranwell's levelled 1951–53 readings provide the comparison. Set against them, the open slacks he measured stand today within about a decimetre of their mid-century level once the climate of the two periods is allowed for (−0.09 ± 0.10 m at the six sites beyond the coastal reach, over sixty-six years; Hollingham, 2026d, §5.7.9). The decline this record resolves is therefore an excursion within a near-stationary century, not the visible part of a longer trend.

The deterioration is concentrated in the summer months. This is consistent with rising evaporative demand under the documented post-2013 summer warming (Figure 3), acting through the atmospheric-draw term as elevated potential evapotranspiration. That route is, if anything, understated by the forcing. Thornthwaite carries the heat index in its denominator and constructs it from the same temperatures, so the formula partly cancels its own response to warming. Across the station record, a 7.2% rise in mean annual temperature (0.73 °C between 1931–1960 and 1996–2025) produced a rise of only 3.0% in mean annual PET, an

elasticity of 0.42. The evaporative demand implied by the temperature record therefore exceeds the demand the model was forced with.

A reduction in effective recharge is also plausible. One example is a shift in the within-month distribution of rainfall, expressed as a declining recharge sensitivity. The two routes sit differently within the framework. The warming signal enters directly through the model's potential-evapotranspiration forcing. A sub-monthly rainfall redistribution, by contrast, is invisible to a monthly model and would surface only as an apparent decline in recharge sensitivity. The framework makes the network-wide signal visible. On the present analysis, it does not resolve which route dominates, or what share of the observed decline either accounts for. Both are noted as forcings the characterization exposes, not results the paper establishes.

5.6 Predictive skill and transferability

The practical contribution of the framework rests on the benchmarking result. In one-step diagnostic mode the state-space and traditional linear models are nearly indistinguishable, because each prediction is corrected by the previous observation. The distinction emerges under iterative simulation. There the SSM achieves a median Nash–Sutcliffe efficiency of 0.77 against the linear model's 0.42. It returns a positive iterative NSE at all but one of the 66 reference wells, against 55 for the linear model. The explicit head-dependent drainage term captures the storage decay that prevents the cumulative drift to which the linear model is prone.

That iterative skill holds far outside the calibration window. The per-well model was driven by the station record from 1930, with the 2005–2026 coefficients and no refitting. It reproduces the seasonal water-table cycle recorded at three levelled pipes in 1951–53 (Ranwell, 1959), with correlations of 0.74–0.91 and the seasonal range to within a tenth of a metre. Those pipes were read fifty-four years before the fitted record begins (Hollingham, 2026d, §5.7.9).

The spatial pattern of the gain is itself a diagnostic. The largest improvements concentrate in the lake-proximal eastern aquifer, where head-dependent drainage is strongest. The smallest occur in the forest interior, where weak drainage makes the storage term nearly dispensable. The improvement map therefore indicates which parts of the aquifer are governed by head-dependent drainage and which approximate a linear climate response. It is a structural characterization obtained without any invasive measurement.

This is the basis for transferability. The framework requires only three fitted parameters per well, a uniform drainage datum, and monthly rainfall and temperature from a single climate station. It does not require hydraulic-conductivity measurements, aquifer-thickness surfaces or lateral-flow modelling. These are precisely the parameter-sparsity constraints that produced calibration difficulties for an earlier distributed-model effort at this site (Betson et al., 2002).

The displacement formulation is the methodological choice that enables this. It references drainage to a datum below ground rather than to the ground surface. In doing so it captures the head-dependence of the drainage response without calibrating the drainage geometry. It also resolves a positive drainage coefficient even in the low-gradient forest interior. The datum

itself is site-specific. The shallowest datum at which the drainage term is identified, and the depth beyond which the drainage flux settles, both follow the site's water-table depths. The datum sweep is therefore repeated at each new site rather than the 3.7 m value being reused (Hollingham, 2026c). With that step, any coastal dune aquifer with a manual dipwell network and a nearby climate station can be parameterized in the same way. The diagnostic syntheses demonstrated here (the t½/ΔNSE/Sy characterization, the coefficient atlas, the residual field) follow directly from that parameterization.

One supplementary step reaches outside this core. Rendering the forest-interception drawdown as a physical reach in metres (Figure 18) requires a hydraulic diffusivity. For this, a single literature conductivity (K ≈ 6 m/day; Betson et al., 2002) and a nominal saturated thickness are assumed, the latter uncertain by a factor of two or more. These enter only that illustrative reach length. The core parameterization, the diagnostic syntheses above and the empirically fitted coastal reach (Figure 19) all remain free of conductivity and thickness inputs. The forest reach is accordingly an order-of-magnitude diagnostic rather than a calibrated prediction.

5.7 Framework limitations

The framework's strengths and its limits are two sides of the same design choice. It fits a lumped mass balance independently at each well and interpolates the results geometrically, so it does not resolve lateral groundwater flow. The interpolated coefficient and head surfaces are aggregators of point responses, not a physical flow model. They cannot represent cluster-to-cluster hydraulic coupling, the lateral redistribution of localized recharge perturbations, or the diffusive response of the aquifer to a change in boundary flux. The affinity analysis and the Darcy field show that the clusters are connected along a continuous downslope drainage pathway. Whether a signal originating in one part of the aquifer propagates to another at magnitudes below the framework's resolution therefore cannot be answered here. That is a question for a calibrated continuous-flow model.

The critical unknowns for such a model are the very quantities the present framework avoids. One is hydraulic conductivity, currently constrained by a single tracer test (6.0 m/day; Betson et al., 2002). The other is aquifer thickness, inferred from a handful of borehole logs across the site. A measured K would also tighten the only K-conditioned output of the present framework, the forest-drawdown reach length in Figure 18. As noted in Section 5.6, however, all other framework outputs (SSM coefficients, WTF specific yields, substrate-gradient interpretation, coastal-retreat regression and residual field) are K-independent and would not change with a measured K. The priority field measurements to enable a flow model are slug tests at representative wells per cluster and a ground-penetrating-radar survey of aquifer thickness. These would also support an independent vadose-zone test of the canopy-driven recharge-suppression signal.

The same campaign would test the structural interpretation of the C3 transition (Section 5.1). A south-west to north-east transect would resolve the depth to the estuarine-clay base, carrying the Caernarfon Bay coastal-retreat rate as a covariate. It would establish whether the eastward-

to-south-westward buffering gradient reflects a single dipping clay base (one thickening sand column) rather than two distinct substrate blocks. It would also separate that structural control from the confounded coastal and plantation-margin influences.

Several further limitations qualify specific results. The monthly measurement resolution cannot isolate individual storm events. The water-table-fluctuation specific yields therefore conflate gravity drainage with capillary-fringe release, and are upper bounds rather than point estimates (Healy and Cook, 2002; Scanlon et al., 2002). Thornthwaite PET is a temperature-and-daylength index whose absolute magnitude is uncertain. The framework's internal consistency does not depend on it, but absolute flux estimates derived from it do.

The coastal-retreat regression is a steady-state spatial model. It resolves the first-order distance gradient but not the multi-year transient over which an erosion signal propagates inland. Direct measurement of present-day retreat rates at the south-western dune front remains a monitoring priority. The water-balance residual shows no spatial structure at the wells, but its monthly series carries a reproducible winter-phased seasonal pattern. As discussed, the attribution of that pattern between a genuine lateral flux and a modelling artefact is unresolved.

Finally, the water table follows the ground closely only once the DEM is smoothed to the scale of the dune ridges ($r = 0.92$ at 300 m, against 0.76 at native resolution). Topography is therefore a proxy for the regional form of the water table, co-determined with it by the underlying till and bedrock architecture. It is not a proxy for local flow paths, and DEM-derived surface routing is not a guide to groundwater flow direction.

6. Conclusions

This study sets out a parameter-sparse state-space framework for characterizing the hydrogeological architecture of a coastal dune aquifer from routine monitoring data. It demonstrates the framework across 88 dipwells and 21 years of monthly records at Newborough Warren. The principal conclusions are:

1. **Five hydrogeological clusters govern site hydrology.** The Pearson affinity audit of hydrograph similarity independently supports cluster assignment for 91% (60 of 66) of reference-network wells (Section 4.3). The cluster-level contrasts in SSM coefficients, drainage half-life and specific yield provide mechanistic substantiation. The five units resolve as a responsive eastern block (C1 Lake Edge, C2 Dune), a buffered Western Residual (C3), and two forested clusters distinguished by topographic position and inferred substrate. C4 Main Forest, on the elevated ridge flank, has high atmospheric draw and weak drainage. C5 Coastal Forest lies on the lower coastal ground. Within the forest zone, elevation alone explains 98% of the variance in atmospheric draw, which identifies position and elevation, rather than canopy, as the dominant within-forest control.

2. **The aquifer's storage varies smoothly, without discontinuity at the cluster margins; the cluster contrasts are surface-mediated.** Once corrected for canopy interception, the C4 forest specific yield falls within the open-dune range. The interpolated specific-yield surface shows no discontinuity at the cluster margins. The C5 estimate sits at the upper edge of that range and is the least reliable in the network, so it is treated as weakly corroborative. The forest clusters' distinctive coefficient signature therefore derives largely from a surface boundary condition, canopy interception suppressing recharge, rather than from a different subsurface architecture.

3. **The displacement-formulation SSM closes the water balance and substantially improves forecast skill.** The three-term model closes the monthly water balance to within 1.8% of total losses at every cluster. This is a property of the fit rather than independent evidence, and the split between drainage and evapotranspiration depends on the datum. The model achieves a median iterative Nash–Sutcliffe efficiency of 0.77 against 0.42 for a traditional linear model, with a positive iterative NSE at 65 of 66 wells against 55 of 66. The spatial pattern of the improvement is itself a non-invasive diagnostic of where head-dependent drainage governs the aquifer.

4. **The water balance closes on average and its residual shows no spatial structure at the wells.** Residuals lie within ±0.01 m/month at 65 of the 66 reference wells and carry no gradient across the site. No well exceeds +0.02 m/month, and the largest positive residuals fall in the open dune and at the lake edge rather than in the forest. No unmodelled lateral flux is required to close the balance at any cluster. A cross-correlation lag test and a seasonal climatology test on the monthly residual likewise return null results for ridge-derived input. The monthly sampling interval is, however, coarse relative to the transit times expected of fracture flow, so the nulls bound what this record can resolve rather than excluding the mechanism. The residual field is a structural diagnostic of model adequacy, not a quantified flux map.

5. **Coastal retreat is a distinct, spatially structured forcing on long-term water-table trends.** A network-scale gradient regression resolves a coast-edge deepening of roughly 26 mm/yr at 150 m from the shoreline, declining over an inland reach of order 1 km. This deepening accounts for about half of the observed summer-minimum decline at the coast-proximal C5 cluster. It also accounts for most of the smaller decline at the Western Residual, half of whose wells lie within the reach. The gradient is corroborated by the progressive deepening of the coast-to-inland head difference along a shore-normal transect. It is spatially structured, strongest at the shore and decaying to nothing over that reach, rather than a uniform decline. The exceptional decline of the coast-proximal C5 cluster is consistent with its coastal position and thin saturated thickness. A contribution from its maturing canopy is not excluded and is left open. Distinguishing this coastal forcing from the climate background is necessary for any correct attribution of network-wide summer-minimum trends.

6. **The framework is transferable.** It characterizes aquifer architecture (storage, recharge sensitivity, storage decay and flow structure) from manual dipwell records and a single

climate station. It does so without coring, hydraulic testing or calibrated continuous-flow modelling. Its one site-specific modelling choice, the drainage datum, is set by repeating the datum sweep of Hollingham (2026c) at each site. Its principal limitation is the converse of this strength. As a geometric aggregator of point responses, it cannot resolve cluster-to-cluster hydraulic coupling. That would need a calibrated flow model, and the hydraulic-conductivity and aquifer-thickness measurements it requires.

The companion paper (Hollingham, 2026b) quantifies the hydrological consequences of the dune-scraping and plantation-clearfell interventions undertaken at the site, within a before-after-control-impact framework.

Source data

For review only — every figure and table below traces to a committed file in the project repository (paths relative to outputs/, or to config.py / data/ where noted). This block lets reviewers verify each value against its producing artefact and is to be removed at production.

Item	Committed source
Table 1 — cluster mechanistic characterization	03_state_space_model/ 03_03_cluster_mechanistic_coefficients.csv
Table 2 — mean monthly head-space water balance	16_water_balance/16_water_bal_table.csv
Table 3 — annual volumetric water balance	16_water_balance/16_water_bal_vol_table.csv
Table 4 — specific yield by cluster (WTF method)	17_wtf_specific_yield/17_wtf_01_sy_estimates.csv
Table 5 — per-well SSM coefficient ranges by cluster	07_spatial_coefficients/07_coeff_05_cluster_ranges.csv
Table 6 — drainage decay half-life $t_{½} = \ln(2)/\beta_3$ by cluster	18_wtf_spatial/18_wtf_05_storage_drainage_index.csv
Table 7 — within-forest spatial predictors of SSM coefficients	10c_forest_zone_analysis/10c_forest_zone_correlations.csv
Table 8 — decomposition of cluster summer-minimum decline	25_coastal_gradient/25_03_cluster_partition.csv
Figure 1 — site topography, geology and monitoring network	01_locations.csv (+ DEM/KML in data/geo/)
Figure 2 — climate forcing over the monitoring period	01_climate.csv
Figure 3 — summer-temperature record	00_climate_summary/00_03_summer_warming_stats.csv
Figure 4 — Ward's dendrogram	02_clustering/02_07_cluster_membership_k5.csv
Figure 5 — spatial cluster distribution (Pearson affinity)	05_pear_membership_audit.csv (+ 01_locations.csv)
Figure 6 — cluster-centroid hydrographs and amplitudes	02_clustering/02_09_cluster_amplitude_summary.csv
Figure 7 — water-balance decomposition by cluster	16_water_balance/16_water_bal_table.csv
Figure 8 — WTF specific-yield surface	18_wtf_spatial/18_wtf_01_well_sy_estimates.csv
Figure 9 — SSM-over-TLM Nash–Sutcliffe improvement	08_model_benchmarking/08_perwell_nse.csv
Figure 10 — optimal drainage datum and R^2 gain	03_state_space_model/03_09_well_optimal_datums.csv; 03_state_space_model/03_08_datum_sensitivity.csv

Figure 11 — β_1 recharge-sensitivity surface	07_spatial_coefficients/07_coeff_maps_data.csv
Figure 12 — β_2 atmospheric-draw surface	07_spatial_coefficients/07_coeff_maps_data.csv
Figure 13 — β_3 drainage-rate surface	07_spatial_coefficients/07_coeff_maps_data.csv
Figure 14 — drainage decay half-life t½	18_wtf_spatial/18_wtf_05_storage_drainage_index.csv
Figure 15 — kriged water table and flow arrows	01b_water_table/01b_01_water_table_flow.png (kriged from the dipwells; 01b_report_numbers.csv)
Figure 16 — SSM water-balance residual at the wells	20_spatial_figures/20_residual_perwell.csv
Figure 17 — coastal-retreat gradient (a, b, c)	25_coastal_gradient/25_report_numbers.csv; 25_coastal_gradient/25_03_cluster_partition.csv; 38_coastal_transect/38_results.txt
Figure 18 — illustrative forest reach	20_spatial_figures/20_drawdown_propagation_nohead.png; schematic, parameters from config.py (FOREST_INTERCEPTION, DRAWDOWN_H0/K/B), rendered via Script 20
Figure 19 — episodic coastal reach	schematic; parameters from 25_coastal_gradient/25_report_numbers.csv and config.py (COAST_RETREAT_*)
Figure 20 — conceptual reach of coastal retreat and the far-field term	09_scraping_intervention/09g_coastal_vs_climate_reach.png; schematic, δ_0/L_cg from 25_coastal_gradient/25_01_panel_fit_parameters.csv, rendered via Script 09g

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author used Anthropic's Claude (successive Claude Opus and Claude Fable models, accessed through the Claude desktop application) for substantial drafting and editing of the manuscript text: composing prose from the author's analytical outputs and decisions, structuring and restructuring sections, and correcting the text against the study's committed analytical outputs. The conception and design of the study, the twenty-one-year manual dipwell record on which it rests, the analytical decisions, and the interpretation of the results are the author's own. Numerical values in the manuscript were verified against the committed outputs of the analysis pipeline, and a decision log records the methodological and editorial calls made by the author during the work. After using this tool the author reviewed and edited the content as needed and takes full responsibility for the content of the published article.

References

Akaike, H. (1974) A new look at the statistical model identification. *IEEE Transactions on Automatic Control* 19(6), 716–723. doi:10.1109/TAC.1974.1100705

Bear, J. (1972) *Dynamics of Fluids in Porous Media*. New York: American Elsevier.

Betson, M., Connell, L. and Bristow, C. (2002) Groundwater Modelling of Newborough Warren. In: Bristow, C. (ed.) *The Impact of Forestry on Coastal Geomorphology at Newborough Warren/Ynys Llanddwyn NNR, SSSI, pSAC. Volume 5: Hydrogeology*. Report to the Countryside Council for Wales, Contract FC 73-05-18. Department of Geological Sciences, UCL, and School of Earth Sciences, Birkbeck, University of London, July 2002.

Curreli, A., Wallace, H., Freeman, C., Hollingham, M., Stratford, C., Johnson, H. and Jones, L. (2013) Eco-hydrological requirements of dune slack vegetation and the implications of climate change. *Science of the Total Environment* 443, 910–919. doi:10.1016/j.scitotenv.2012.11.035

Davy, A.J., Grootjans, A.P., Hiscock, K. and Petersen, J. (2006) *Development of Eco-Hydrological Guidelines for Dune Habitats — Phase 1*. English Nature Research Reports No. 696. Peterborough: English Nature.

Desbarats, A.J., Logan, C.E., Hinton, M.J. and Sharpe, D.R. (2002) On the kriging of water table elevations using collateral information from a digital elevation model. *Journal of Hydrology* 255(1–4), 25–38. doi:10.1016/S0022-1694(01)00504-2

Fetter, C.W. (2001) *Applied Hydrogeology* (4th ed.). Upper Saddle River, NJ: Prentice Hall.

Forgrave, A. (2020) Dramatic pictures show the beautiful little corner of North Wales that's fast being lost to the sea. *North Wales Live / Daily Post*, 22 January 2020.

Freeman, S. (2008) *The hydrological impact of Corsican pine plantation on the dune system at Newborough Warren, Anglesey*. Unpublished BSc Honours dissertation, University of Wales, Bangor.

Freeze, R.A. and Cherry, J.A. (1979) *Groundwater*. Englewood Cliffs, NJ: Prentice-Hall.

Gerla, P.J. (2019) Monitoring and modeling the effect of agricultural drainage and recent channel incision on adjacent groundwater-dependent ecosystems. *Water* 11(4), 863. doi:10.3390/w11040863

Healy, R.W. and Cook, P.G. (2002) Using groundwater levels to estimate recharge. *Hydrogeology Journal* 10(1), 91–109. doi:10.1007/s10040-001-0178-0

Hollingham, M. (2026b) Hydrological response of a coastal dune aquifer to plantation clearfell and dune scraping: a BACI analysis of a 21-year manual record. *In preparation*.

Hollingham, M. (2026c) Intercept, datum and record: what a linear-reservoir model fitted to monthly groundwater levels can tell you. *In preparation*.

Hollingham, M. (2026d) Hydrogeological Dynamics, Behavioural Clustering and Management Intervention Analysis at Newborough Warren Coastal Sand Dune Aquifer, Wales. Technical report, October 2026. https://newbroman.github.io/Newborough_Hydrology/

Knotters, M. and van Walsum, P.E.V. (1997) Estimating fluctuation quantities from time series of water-table depths using models with a stochastic component. *Journal of Hydrology* 197(1–4), 25–46. doi:10.1016/S0022-1694(96)03278-7

Pye, K. and Blott, S.J. (2024) *Dune rejuvenation AfterLIFE options for Twyni Penrhos, Newborough.* NRW Evidence Report No. 787, 73 pp. Bangor: Natural Resources Wales.

Ranwell, D.S. (1958) Movement of vegetated sand dunes at Newborough Warren, Anglesey. *Journal of Ecology* 46, 83–100.

Ranwell, D.S. (1959) Newborough Warren, Anglesey. I. The dune system and dune slack habitat. *Journal of Ecology* 47, 571–601.

Rao, A.R. and Srinivas, V.V. (2006) Regionalization of watersheds by fuzzy cluster analysis. *Journal of Hydrology* 318(1–4), 57–79.

Rao, P., Wang, Y., Liu, Y., Wang, X., Hou, Y., Pan, S., Wang, F. and Zhu, D. (2022) A comparison of multiple methods for mapping groundwater levels in the Mu Us Sandy Land, China. *Journal of Hydrology: Regional Studies* 43, 101189. doi:10.1016/j.ejrh.2022.101189

Rhind, P.M., Blackstock, T.H., Hardy, H.S., Jones, R.E. and Sandison, W. (2001) The evolution of Newborough Warren dune system with particular reference to the past four decades. In: Houston, J.A., Edmondson, S.E. and Rooney, P.J. (eds) *Coastal Dune Management.* Liverpool: Liverpool University Press, pp. 345–379.

Robins, N.S., Pye, K. and Wallace, H. (2013) Dynamic coastal dune spit: the impact of morphological change on dune slacks at Whiteford Burrows, South Wales, UK. Journal of Coastal Conservation 17(3), 473–482. doi:10.1007/s11852-013-0245-4

Rousseeuw, P.J. (1987) Silhouettes: a graphical aid to the interpretation and validation of cluster analysis. *Journal of Computational and Applied Mathematics* 20(1), 53–65. doi:10.1016/0377-0427(87)90125-7

Scanlon, B.R., Healy, R.W. and Cook, P.G. (2002) Choosing appropriate techniques for quantifying groundwater recharge. *Hydrogeology Journal* 10(1), 18–39. doi:10.1007/s10040-001-0176-2

Stratford, C., Ratcliffe, J., Hughes, A.G., Roberts, J. and Robins, N.S. (2007) Complex interaction between shallow groundwater and changing woodland, surface water, grazing and other influences in partly wooded duneland in Anglesey, Wales. In: Ribeiro, L., Chambel, A. and Condesso de Melo, M.T. (eds) *IAH 35th Congress, Groundwater and Ecosystems*, Lisbon.

UKHO (2026) *Admiralty Tide Tables, Volume 1B* (NP201B-26). Taunton: United Kingdom Hydrographic Office.

Varouchakis, E.A. and Hristopulos, D.T. (2013) Comparison of stochastic and deterministic methods for mapping groundwater level spatial variability in sparsely monitored basins. *Environmental Monitoring and Assessment* 185(1), 1–19. doi:10.1007/s10661-012-2527-y

Ward, J.H. (1963) Hierarchical grouping to optimize an objective function. *Journal of the American Statistical Association* 58(301), 236–244.

Young, P.C. (2011) *Recursive Estimation and Time-Series Analysis: An Introduction for the Student and Practitioner*. Berlin: Springer.